

Computer Architecture Review





Operating Systems
Wenbo Shen

ENIAC (1946)

■ Electronic Numerical Integrator and Calculator


- Considered by most to be the first electronic computer
- Vacuum tubes, punchcards
- 50,000+ times slower than my laptop
- 100ft long
 - ▶ each register: 20ft
- Programming with wires




Von-Neumann

- In 1944, John von Neumann joined ENIAC
- He wrote a memo about computer architecture, formalizing ENIAC ideas
 - Eckert and Mauchly have pretty much been forgotten (they were in the trenches)
- These ideas became the **Von Neumann architecture model**
 - A processor that performs operations and controls all that happens
 - A **memory** that contains **code and data**
 - **I/O** of some kind


Von-Neumann Model



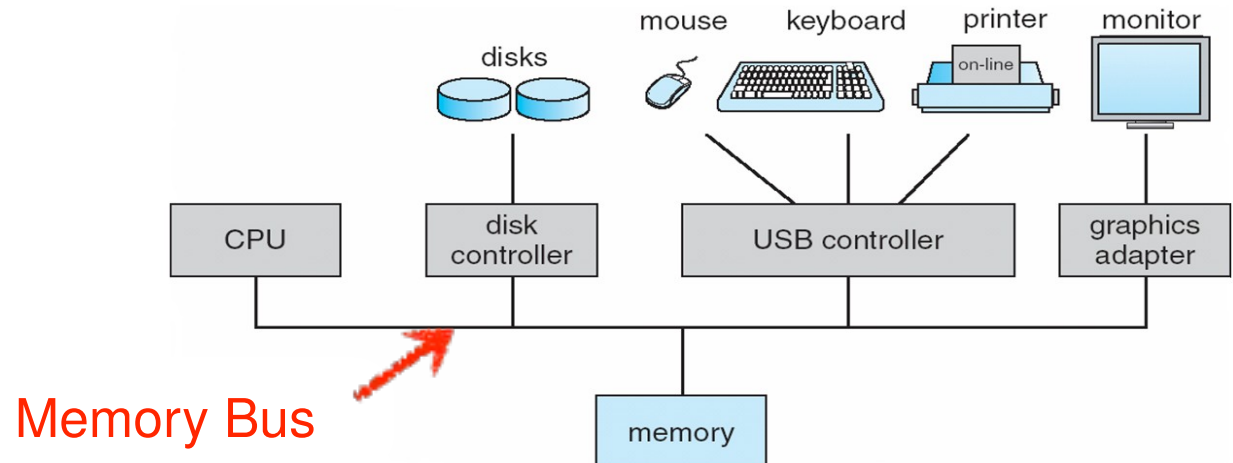
- Amazingly, it's still possible to think of the computer this way at a conceptual level (model from ~70 years ago!!!)
- A computer today has
- some differences



Von-Neumann Model



- Amazingly, it's still possible to think of the computer this way at a conceptual level (model from ~70 years ago!!!)
- A computer today has
- some differences



Data Stored in Memory

- All “information” in the computer is in binary form
 - Since Claude Shannon’s M.S. thesis in the 30’s
 - 0: zero voltage, 1: positive voltage (e.g., 5V)
 - bit: the smallest unit of information (0 or 1)
- The basic unit of memory is a byte
 - 1 Byte = 8 bits, e.g., “0101 1101”
- Each byte in memory is labeled by a unique **address**
- All addresses in the machine have the same number of bits
 - e.g., 64-bit addresses on ARM64
- The processor has instructions that say “Read the byte at address X and give me its value” and “Write some value into the byte at address X”

Conceptual View of Memory

address	content
0000 0000 0000 0000	0110 1110
0000 0000 0000 0001	1111 0100
0000 0000 0000 0010	0000 0000
0000 0000 0000 0011	0000 0000
0000 0000 0000 0100	0101 1110
0000 0000 0000 0101	1010 1101
0000 0000 0000 0110	0000 0001
0000 0000 0000 0111	0100 0000
0000 0000 0000 1000	1111 0101

...

...

Conceptual View of Memory

address	content
0000 0000 0000 0000	0110 1110
0000 0000 0000 0001	1111 0100
0000 0000 0000 0010	0000 0000
0000 0000 0000 0011	0000 0000
0000 0000 0000 0100	0101 1110
0000 0000 0000 0101	0100 0000
0000 0000 0000 0110	1111 0101
0000 0000 0000 0111	...
...	...

**At address 0000 0000 0000 0010
the content is 0000 0000**

Conceptual View of Memory

address	content
0000 0000 0000 0000	0110 1110
0000 0000 0000 0001	1111 0100
0000 0000 0000 0010	0000 0000
0000 0000 0000 0011	0000 0000
0000 0000 0000 0100	0101 1110
0000 0000 0000 0101	
0000 0000 0000 0110	
0000 0000 0000 0111	
0000 0000 0000 1000	1111 0101

At address 0000 0000 0000 0100
the content is 0101 1110

Both Code and Data in Memory

- Once a program is loaded in memory, its address space contains both **code** and **data**
- To the CPU those are not really different, but the programmer knows which bytes are data and which are code
 - Always conveniently hidden from you if you've never written assembly
 - But we'll have to keep code/data straight in these lecture notes

Code

Data

Example Address Space


Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory


We need a CPU

- So now we have a memory in which we can store/retrieve bytes at precise location
- These bytes presumably have some useful meaning to us
 - e.g., integers, ASCII codes of characters, floating points numbers, RGB values
 - e.g., instructions that specify what to do with the data; when you buy a processor, the vendor defines the instruction set (e.g., instruction “0010 1101” means “increment some useful counter”)
- The CPU is the piece of hardware that modifies the content of memory
 - In fact, one can really think of the CPU as a device that takes use from one memory state (i.e, all the stored content) to another memory state (some new, desired stored content)

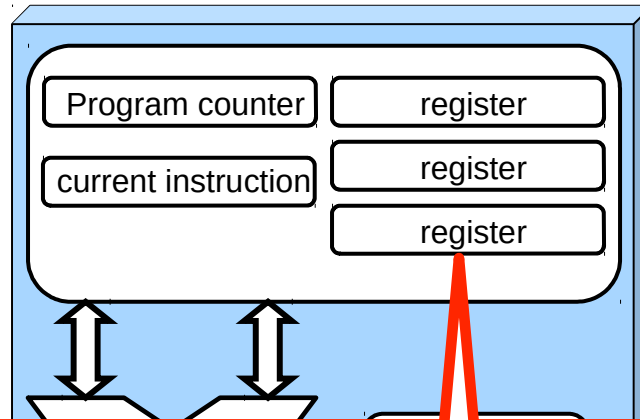
What's in the CPU?



What's in the CPU?



What's in the CPU?



Registers: the “variables” that hardware instructions work with

Data can be loaded from memory into a register

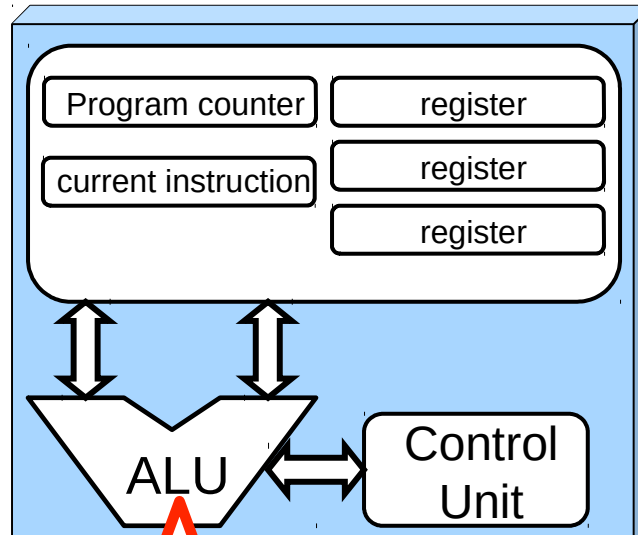
Data can be stored from a register back into memory

Operands and results of computations are in registers

Accessing a register is really fast

There is a limited number of registers

What's in the CPU?

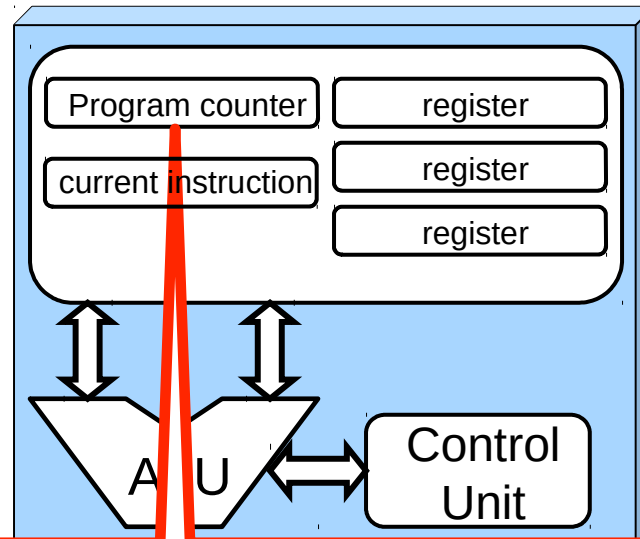


Arithmetic and Logic Unit: what you do computation with

Used to compute a value based on current register values and store the result back into a register

+, *, /, -, OR, AND, XOR, etc.

What's in the CPU?

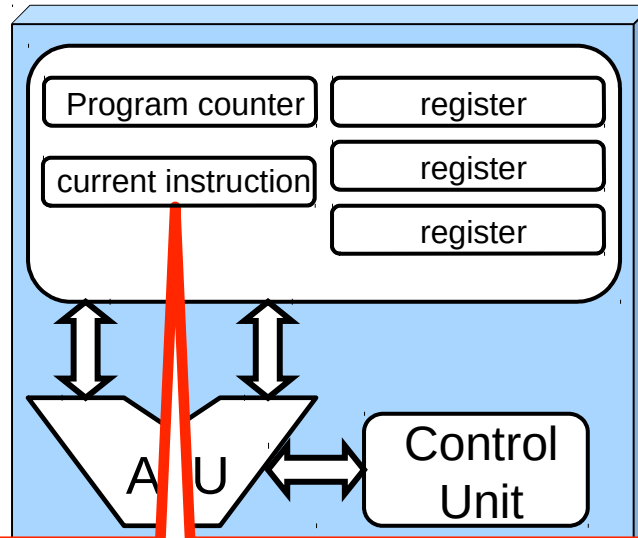


Program Counter: Points to the next instruction

Special register that contains the address in memory of the next instruction that should be executed

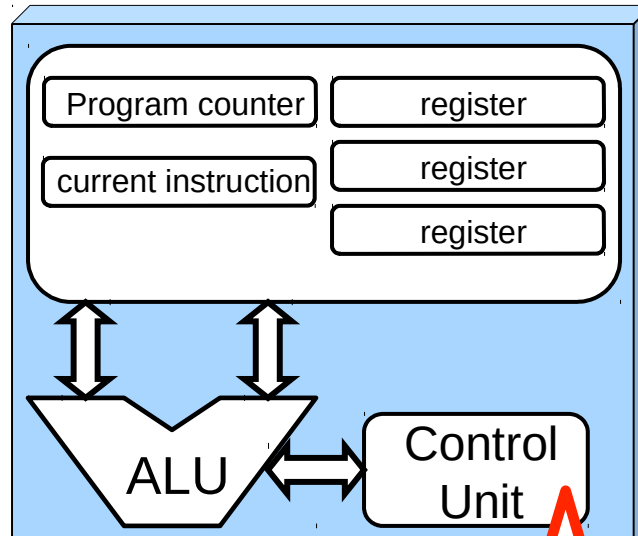
(gets incremented after each instruction, or can be set to whatever value whenever there is a change of control flow)

What's in the CPU?



Current Instruction: Holds the instruction that's currently being executed

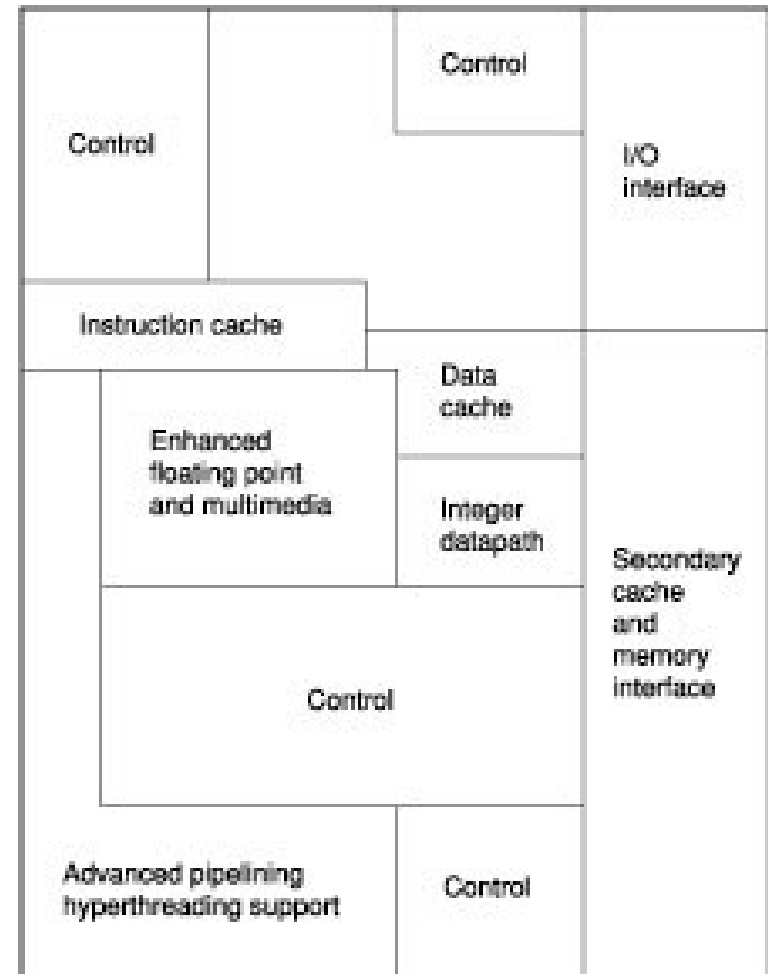

What's in the CPU?



Control Unit: Decodes instructions and make them happen

Logic hardware that decodes instructions (i.e., based on their bits) and sends the appropriate (electrical) signals to hardware components in the CPU

The CPU in its “Glory”



Fetch-Decode-Execute Cycle

■ The Fetch-Decode-Execute cycle

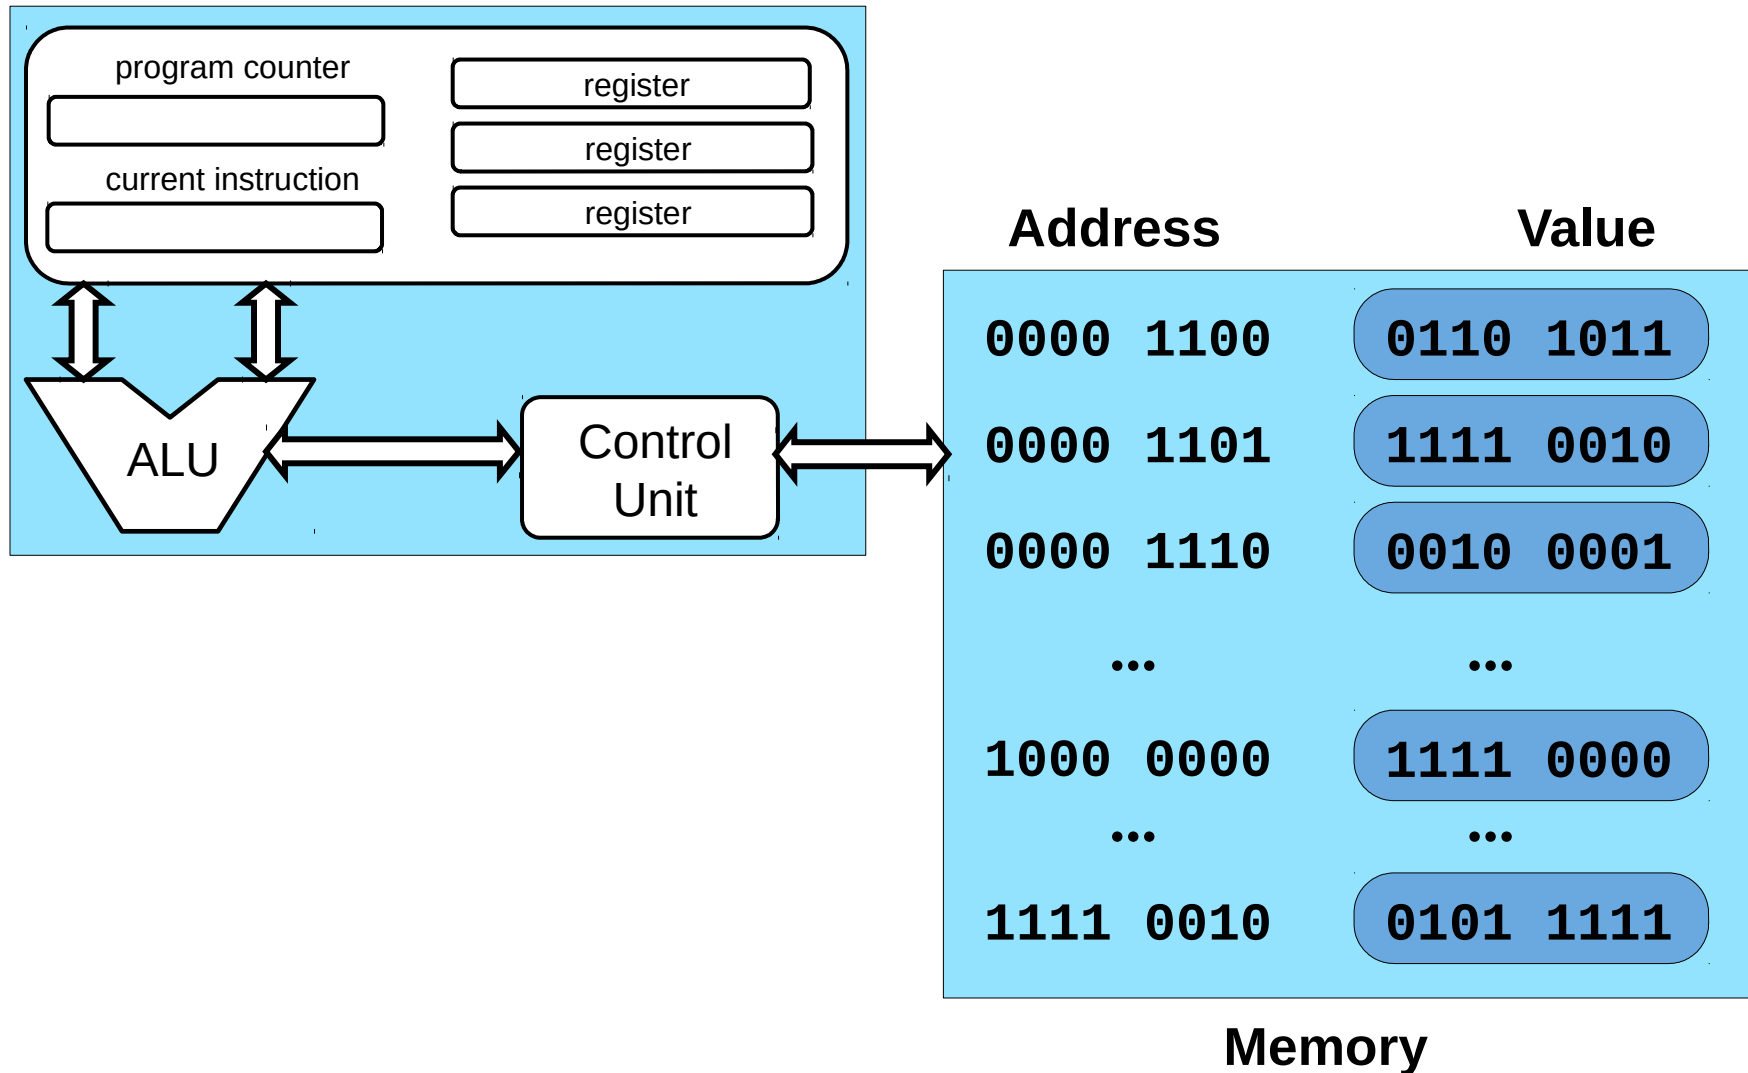
- The control unit fetches the next program instruction from memory
 - ▶ Using the program counter to figure out where that instruction is located in the memory
- The instruction is decoded and signals are send to hardware components
 - ▶ e.g., is the instruction loading something from memory? is it adding two register values together?
- Operands are fetched from memory and put in registers, if needed
- The ALU executes computation, if any, and store results in the registers
- Register values are stored back to memory, if needed
- Repeat

■ Computers today implement MANY variations on this model

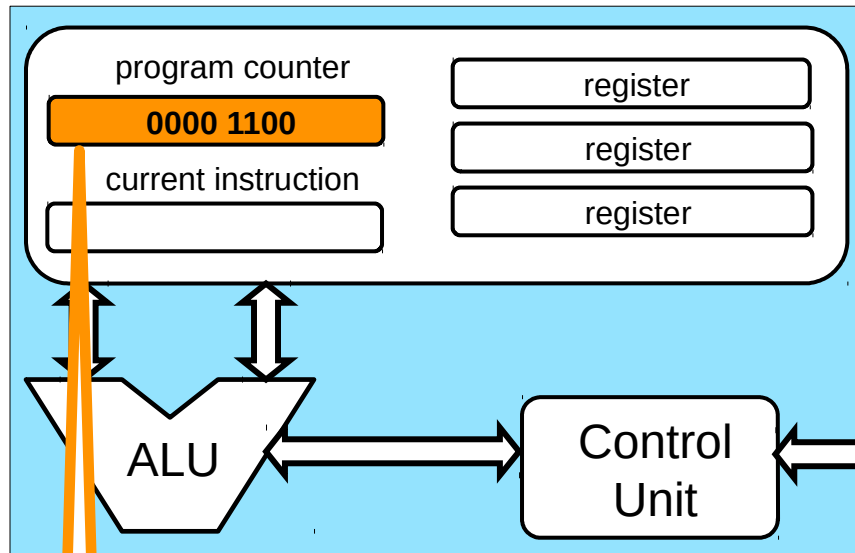
■ But one can still program with the above model in mind

- but certainly without (fully) understanding performance issues

Fetch-Decode-Execute



Fetch-Decode-Execute

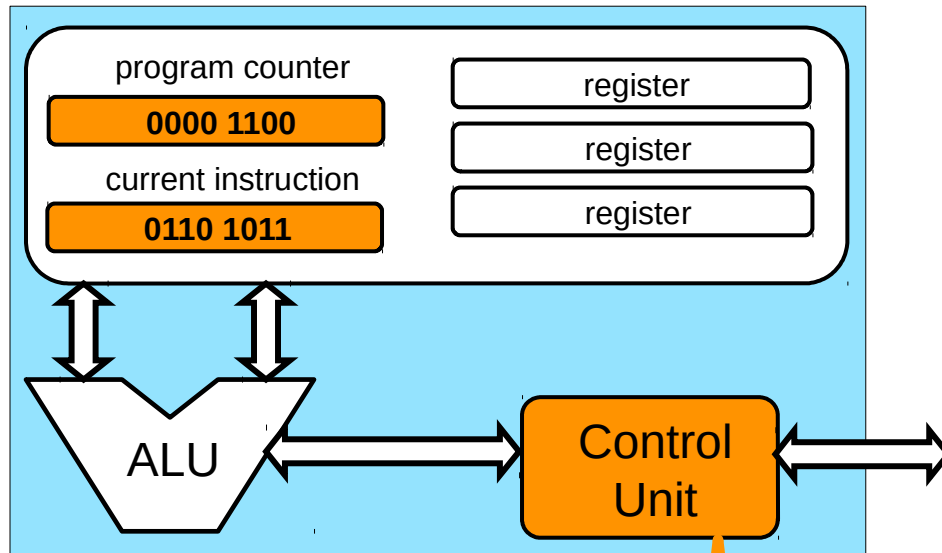


Somehow, the program counter is initialized to some content, which is an address (we'll see how that happens much later)

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

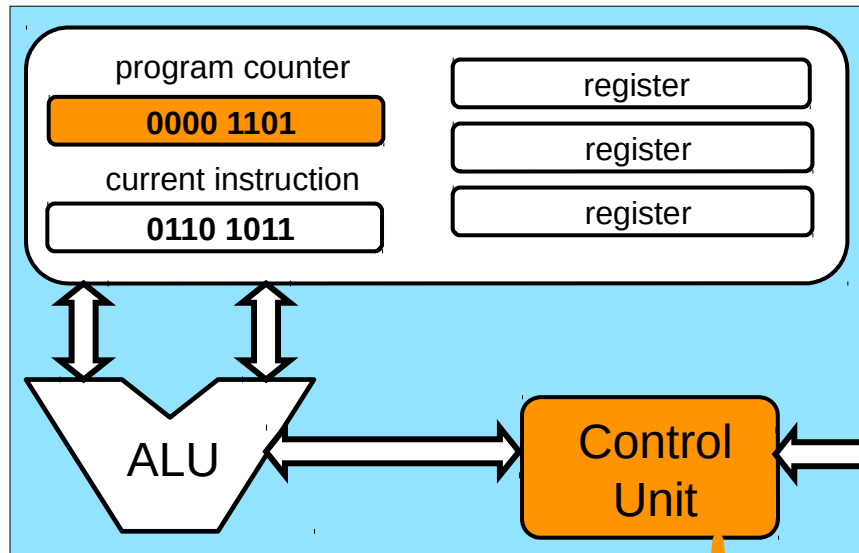


Fetch the content (instruction) at address 0000 1100, which is “0110 1011”, and store it in the “current instruction” register

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

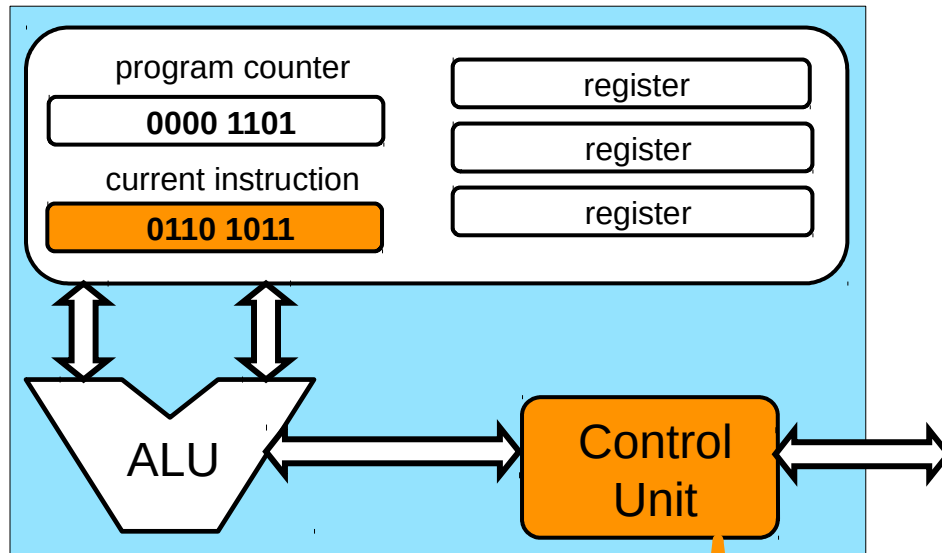


Increment the program counter

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

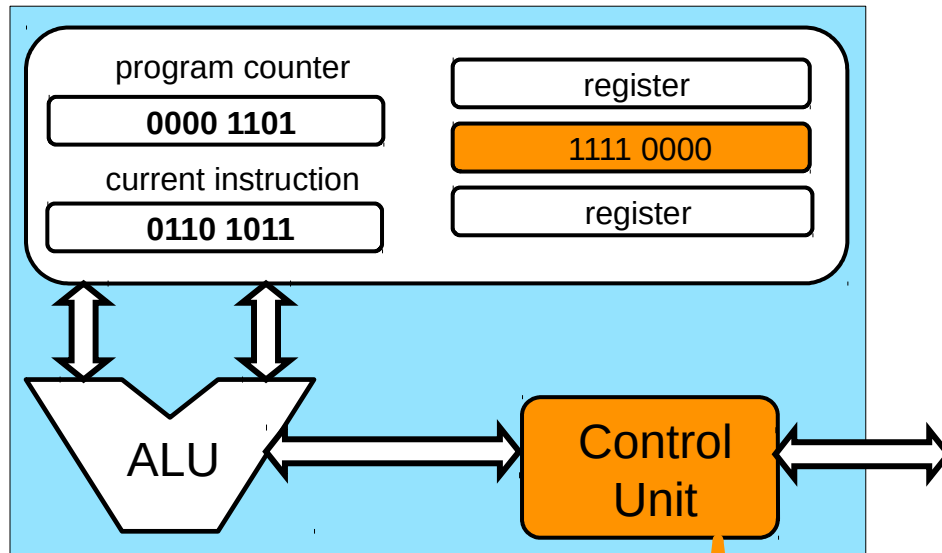


Decode instruction "0110 1011". Say it means: "Load the value at address 1000 0000 and store it in the second register"

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

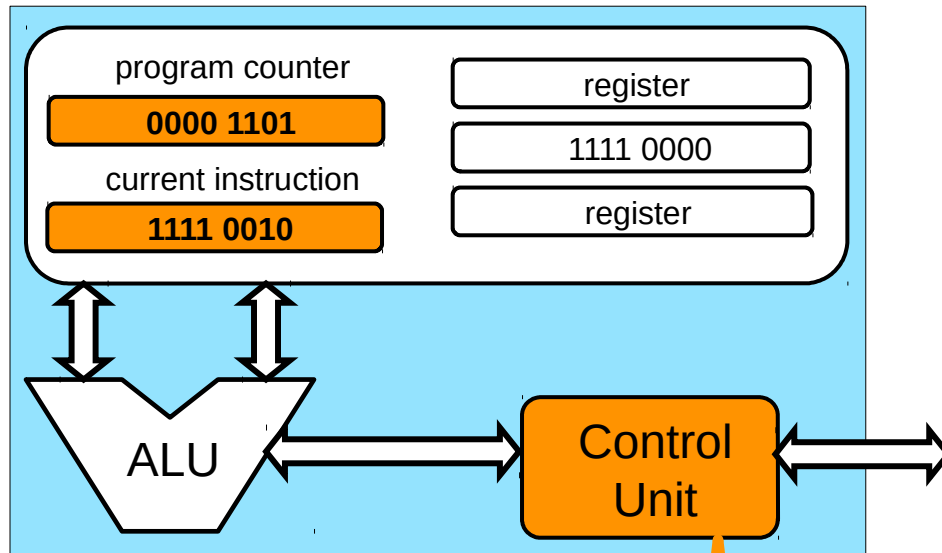


Send signals to all hardware components to **execute** the instruction: load the value at address 1000 0000, which is "1111 0000" and store it in the second register

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

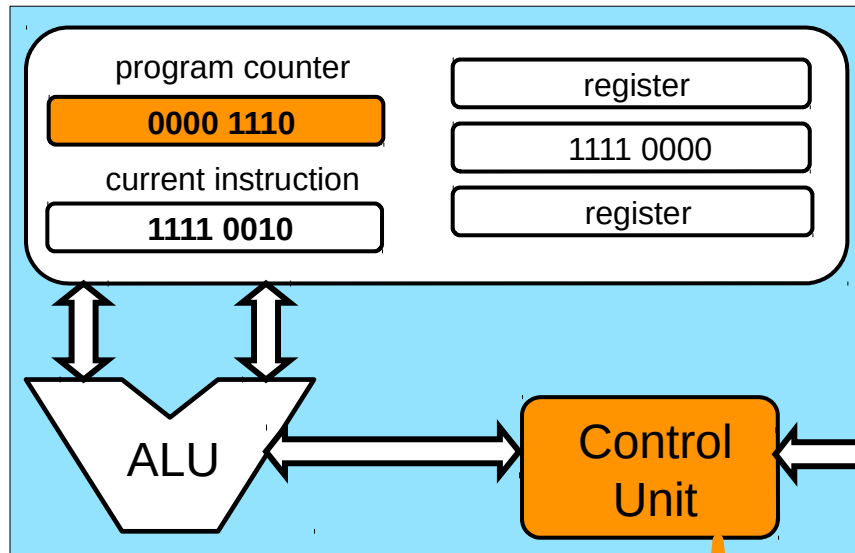


Fetch the content (instruction) at address 0000 1101, which is “1111 0010”, and store it in the “current instruction” register

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

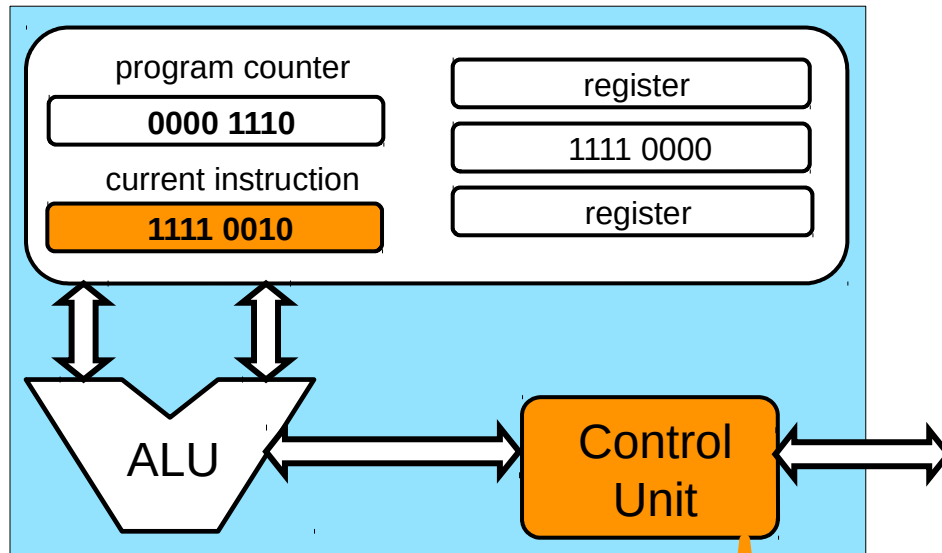


Increment the program counter

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

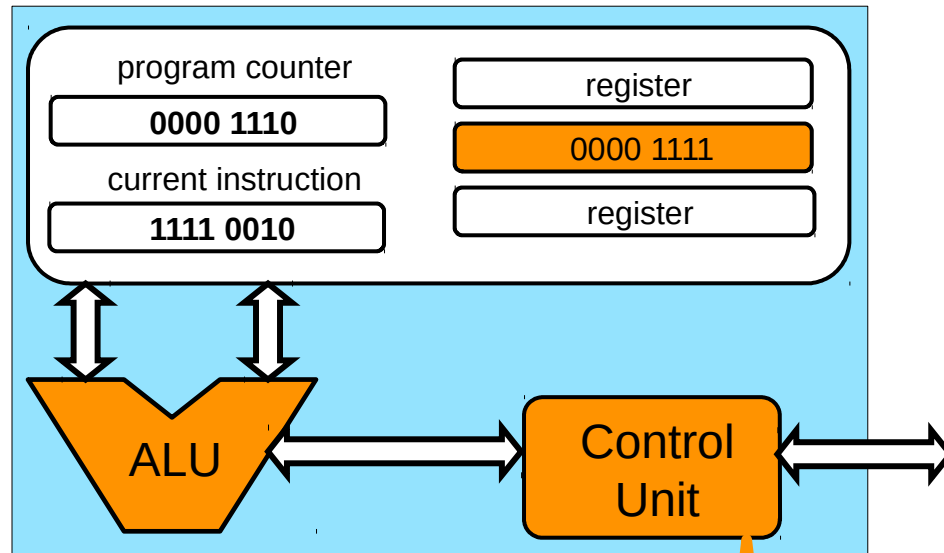


Decode instruction "1111 0010". Say it means: "Do a logical NOT on the second register"

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

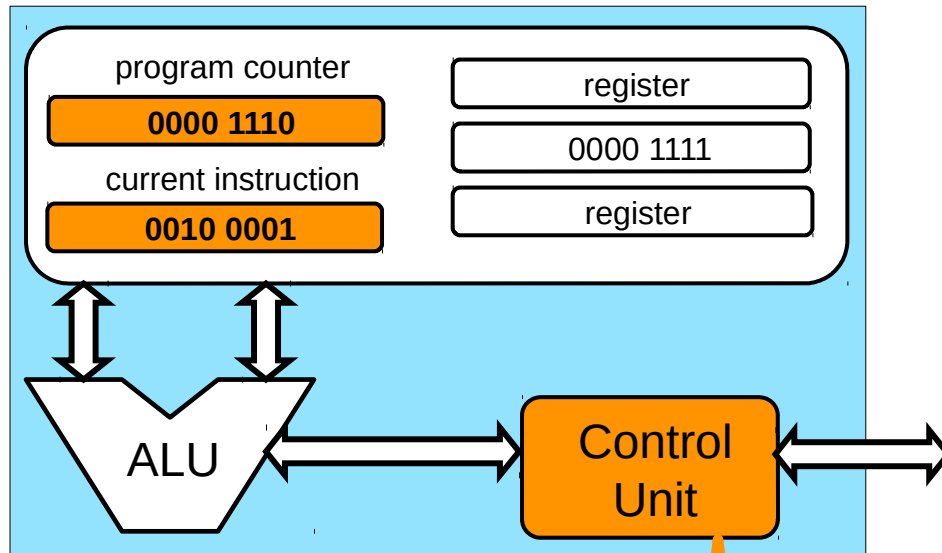


Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Send signals to all hardware components to **execute** the instruction: do a logical NOT on the second register

Fetch-Decode-Execute

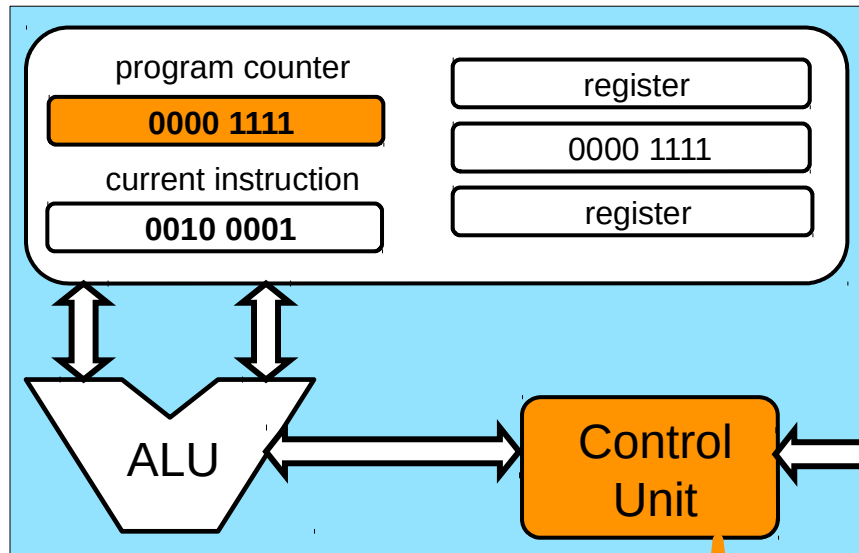


Fetch the content (instruction) at address 0000 1110, which is “0010 0001”, and store it in the “current instruction” register

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

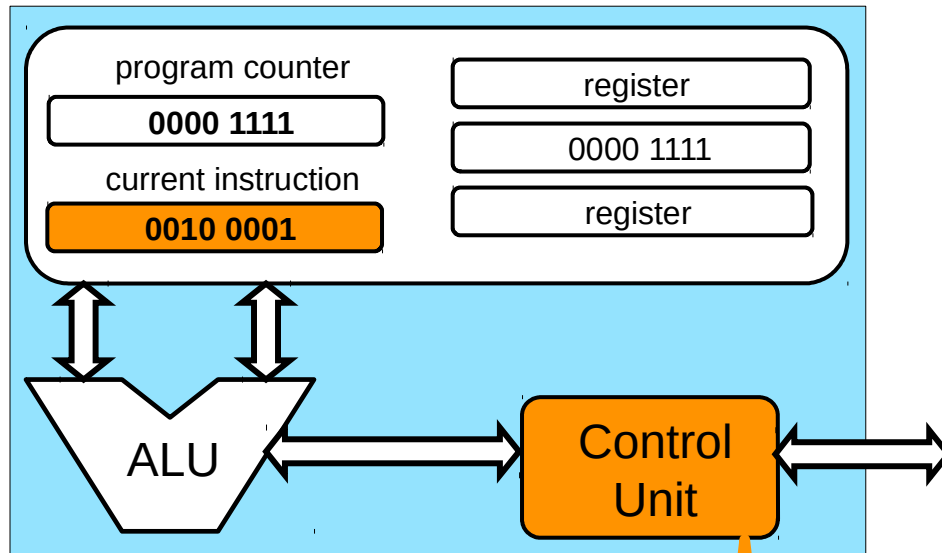


Increment the program counter

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Fetch-Decode-Execute

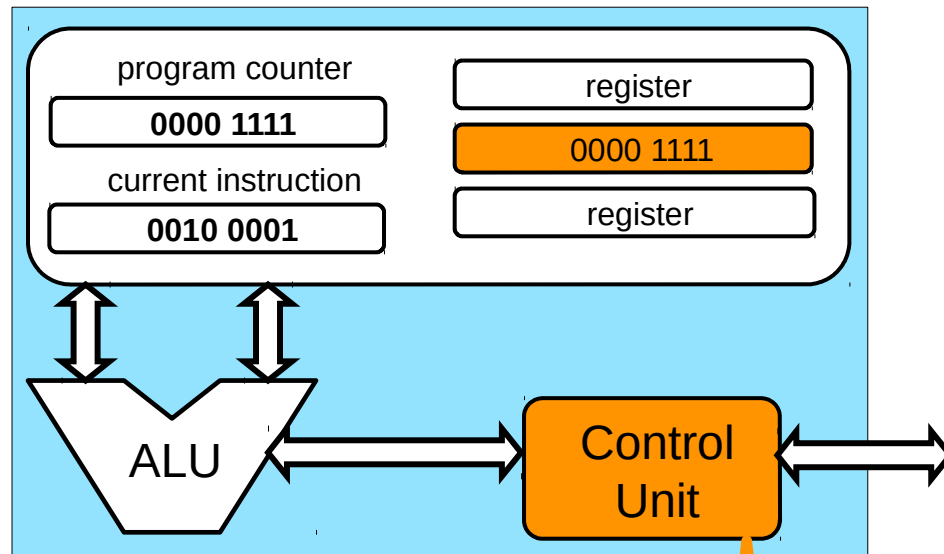


Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111

Memory

Decode instruction "0010 0001". Say it means: "Store the value in the second register to memory at address 1111 0010"

Fetch-Decode-Execute



Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0000 1111

Memory

Send signals to all hardware components to **execute** the instruction: store the value in the second register, which is 0000 1111, to memory at address 1111 0010


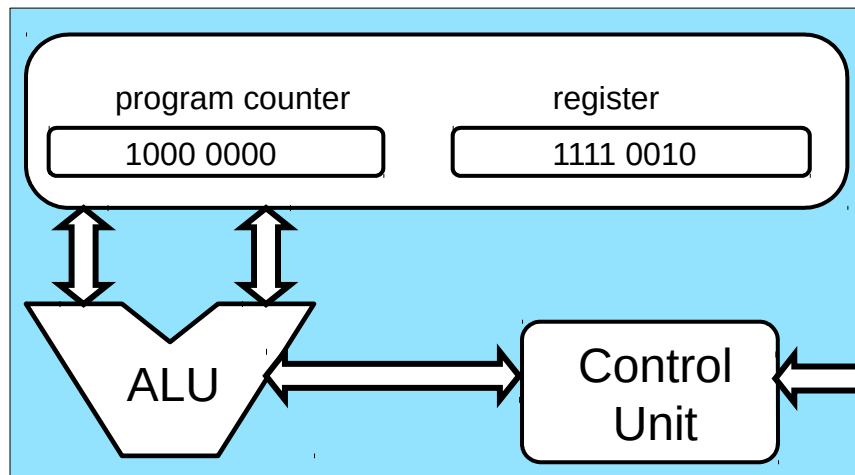
Fetch-Decode-Execute

- This is only a simplified view of the way things work
- The “control unit” is not a single thing
 - Control and data paths are implemented by several complex hardware components
- There are multiple ALUs, there are caches, there are multiple CPUs in fact (“cores”)
- Execution is pipelined: e.g., while one instruction is fetched, another is executed
- Decades of computer architecture research have gone into improving performance, thus often leading to staggering hardware complexity
 - Doing smart things in hardware requires more logic gates and wires, thus increasing processor cost
- But conceptually, fetch-decode-execute is it

In-Class Exercise

- With the following (totally made up and strange, but small) instruction set definition and with this machine state, what is the new memory state after execution completes?


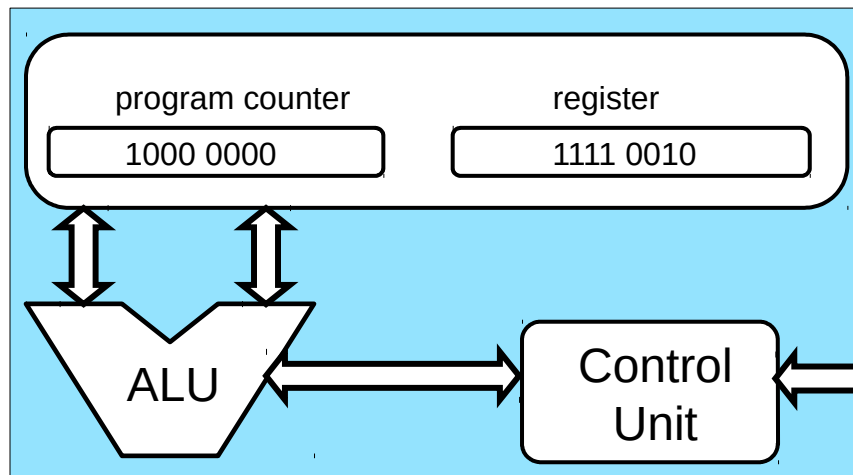
code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address NOT(register)



In-Class Exercise

- Fetch the instruction: “1111 0000”


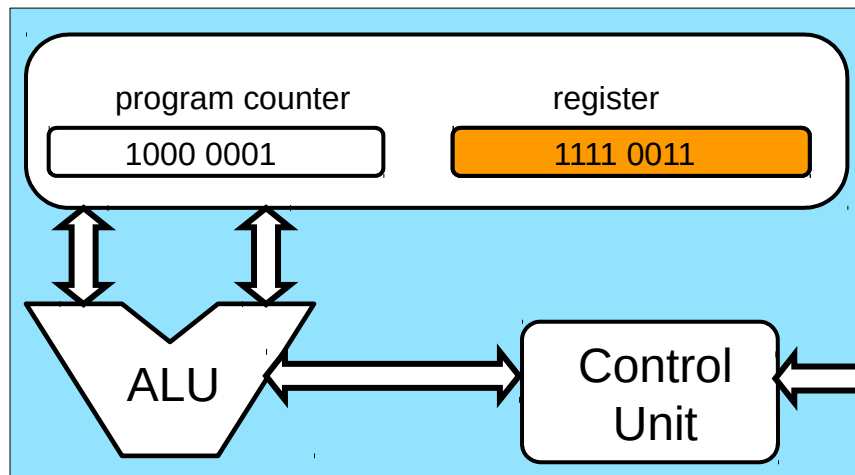
code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address NOT(register)



In-Class Exercise

- Fetch the instruction: “1111 0000”
- Execute it: increment register to value “1111 0011”


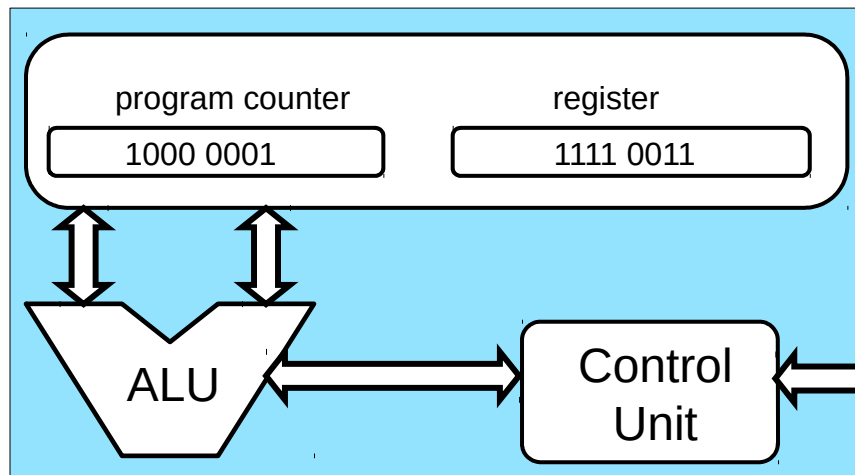
code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address NOT(register)



In-Class Exercise

- Fetch the instruction: “1111 0000”
- Execute it: increment register to value “1111 0011”
- Fetch the next instruction: “0101 1111”


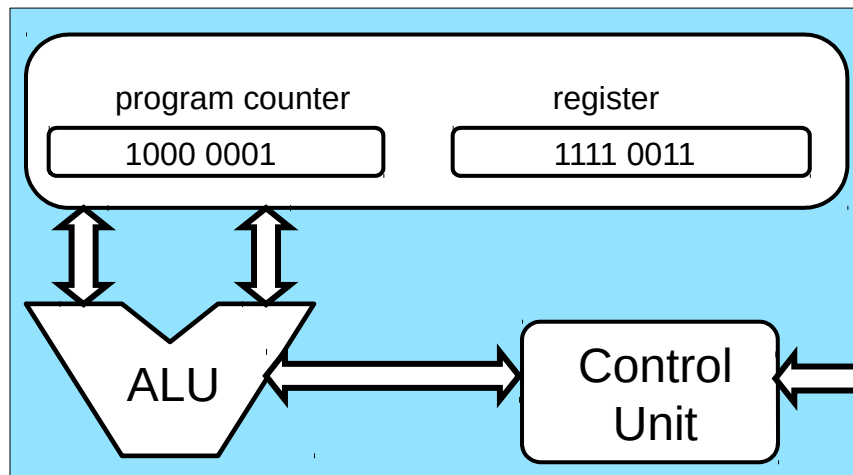
code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address NOT(register)



In-Class Exercise

- Fetch the instruction: “1111 0000”
- Execute it: increment register to value “1111 0011”
- Fetch the next instruction: “0101 1111”
- Execute it: save value “1111 0011” to address “0000 1100”

code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address NOT(register)



Direct Memory Access

- DMA is used in all modern computers
- It's a way for the CPU to let memory-I/O operations (data transfers) occur independently
- Say you want to write 1 GiB from memory to some external device like a disk, network card, graphics card, etc.
- The CPU would be busy during this slow transfer
 - Load from memory into registers, write from registers to disk, continuously
- Instead, a convenient piece of hardware called the DMA controller can make data transfer operations independently of the CPU
- The CPU simply “tells” the DMA controller to initiate a transfer
 - Which is done by writing to some registers of the DMA controller
- When the transfer completes, the DMA controller tells the CPU “it's done” (by generating an interrupt)
 - More on interrupts later
- In the meantime, the CPU can do useful work, e.g., run programs

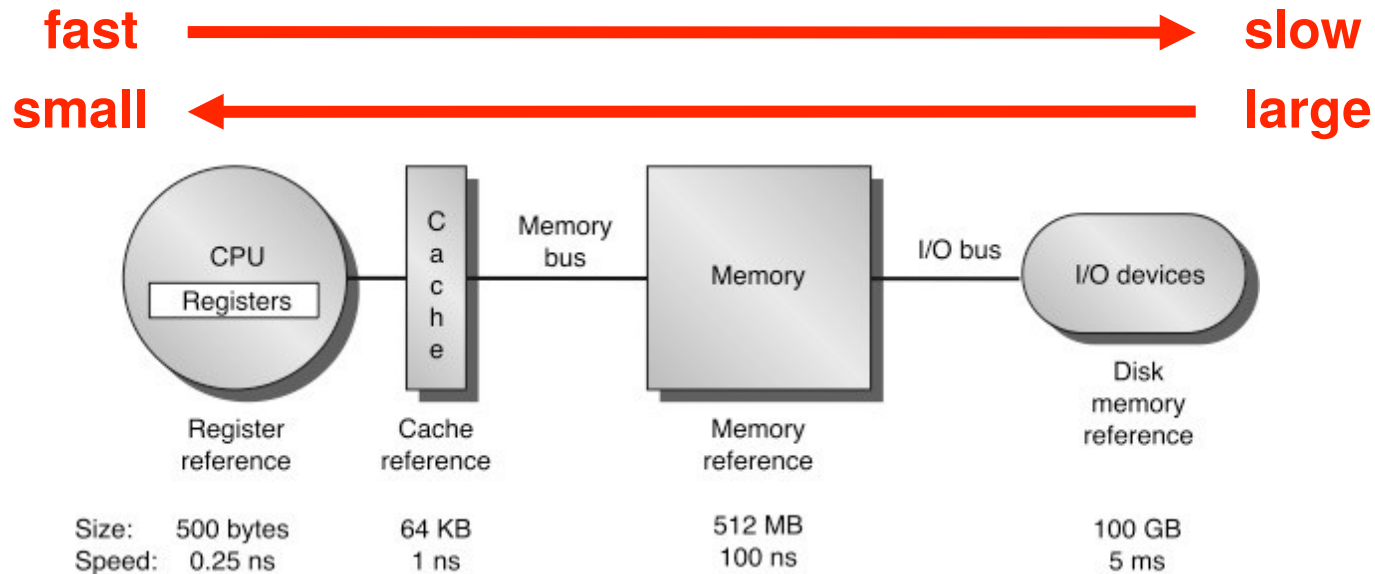
DMA is not completely free

- To perform data transfers the DMA controller uses the memory bus
- In the meantime, the code executed by the CPU likely also uses the memory bus
- Therefore, the two can interfere
- There are several modes in which this interference can be managed
 - DMA has priority
 - CPU has priority
- But in general, using DMA leads to much better performance anyway

Coping with Slow RAM

- Everybody would like to have a computer with a very large and very fast memory
- Unfortunately, technology (affordably) allows for either slow and large or fast and small
- We need large main memories for large programs and data
- Therefore, **from the CPU's perspective, main memory is slooooooow**
- What we do: we play a trick to provide the illusion of a fast memory
- This trick is called the memory hierarchy

The Memory Hierarchy



- Real-world has multiple levels of caches (L1, L2, L3)
- Chunks of data are brought in from far-away memory and are copied and kept around in nearby memory
- Yes, the same data exists in multiple levels of memory at once
- Miss: when a data item is not found in a level (e.g., L1 cache miss)
- Hit: when a data item is found in a level (e.g., L2 cache hit)

Caching

- Whenever your program accesses a byte of memory what happens is:
 - That byte's value is brought from sloooooow memory into the fast cache
 - byte values around the byte you accesses are also brought from sloooooow memory into the fast cache
- Analogy:
 - You need a book from the library
 - You go there and find the book on the many shelves of the library
 - You bring back home all books on that shelf and put them on your own bookshelf in your house
 - Next time you need that book or one of the books “around it”, it will take you no time at all to get it
 - ▶ Presumably all books on a shelf at the library are about the same topic, so you'll need the books around the book you wanted in the first place

Why Does it Work?

- **Temporal Locality**: a program tends to reference addresses it has recently referenced
 - The first access, you pay the cost of going to far-away/slow memory to fetch the counter's content
 - Subsequent accesses are fast
 - This is the “I need that book again” analogy
- **Spatial Locality**: a program tends to reference addresses next to addresses it has recently referenced
 - The first access of array element i may be costly
 - But the first access of array element $i+1$ is fast (in the chunk)
 - This is the “I need another book on that same shelf” analogy


Memory Tech. and Management

Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk and others
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000.000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape

- Main memory and disk are managed by the OS
- When dealing with a “slow” level, it pays off more to try being “clever” (i.e., spending more time trying to make good decisions)
 - Part of why OSeS are doing complicated things, as opposed to hardware which tries to do simple things fast

SMP Systems


■ Symmetric multi processors




Issue: Cache coherency
(see textbook)

Moore's Law

- In 1965, Gordon Moore (co-founder of Intel) predicted that **transistor density** of semiconductor chips would **double** roughly **every 24 months** (often “misquoted” as 18 months)
 - He was right
 - But, the law was often wrongly interpreted as: “Computers get twice as fast every 2 years”
 - This wrong interpretation was true for a while, but no longer...



Moore's Law



This did not happen!!

Multi-core Chips

- Constructors cannot increase clock rate further
 - Power/heat issues
- They bring you **multi-core processors**
 - Multiple “low” clock rate processors on a chip
- It's really a solution to a problem, not a cool new advance
 - Even though there are many cool/interesting things about multicore processors
- Most users/programmers would rather have a 100GHz core than 50 2GHz cores

Multi-Core Systems






Figure 1.7 from the book




More realistic picture

Multi-Proc Multi-Core Systems



Takeaway

- Each byte in memory is labeled by a unique **address**
- CPU architecture
- The **Fetch-Decode-Execute** cycle



Conclusion

- If you want to know more
 - Take a computer architecture course
 - Classic Textbook by Patterson and Hennessy
- Reading assignment:
Sections 1.2 and 1.3

