

Threads



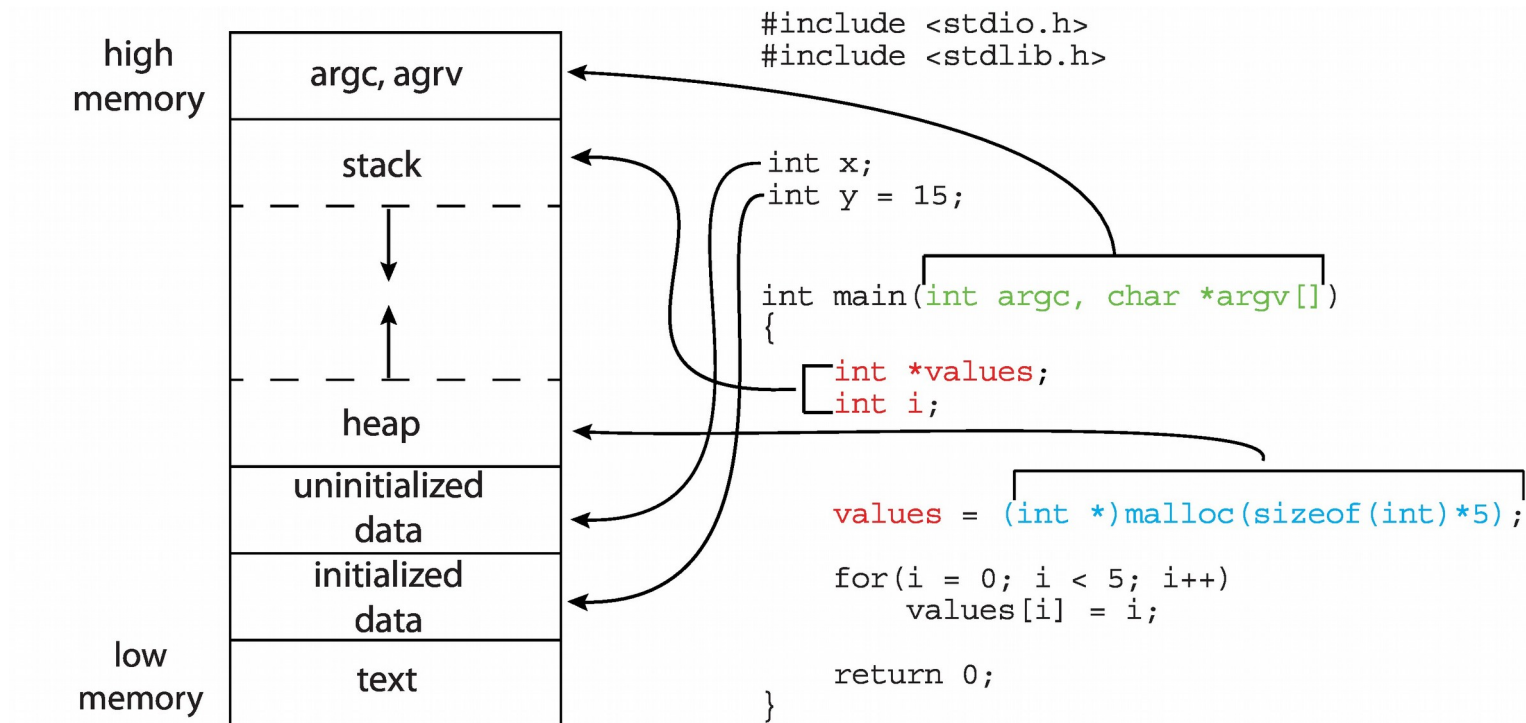
Operating Systems
Wenbo Shen

Revisit - Process Concept

■ Process =

- **code** (also called the **text**)
 - ▶ initially stored on disk in an executable file
- **program counter**
 - ▶ points to the next instruction to execute (i.e., an address in the code)
- content of the processor's **registers**
- a runtime **stack**
- a **data section**
 - ▶ global variables (.bss and .data in x86 assembly)
- a **heap**
 - ▶ for dynamically allocated memory (malloc, new, etc.)


Revisit - Memory Layout of a C Program



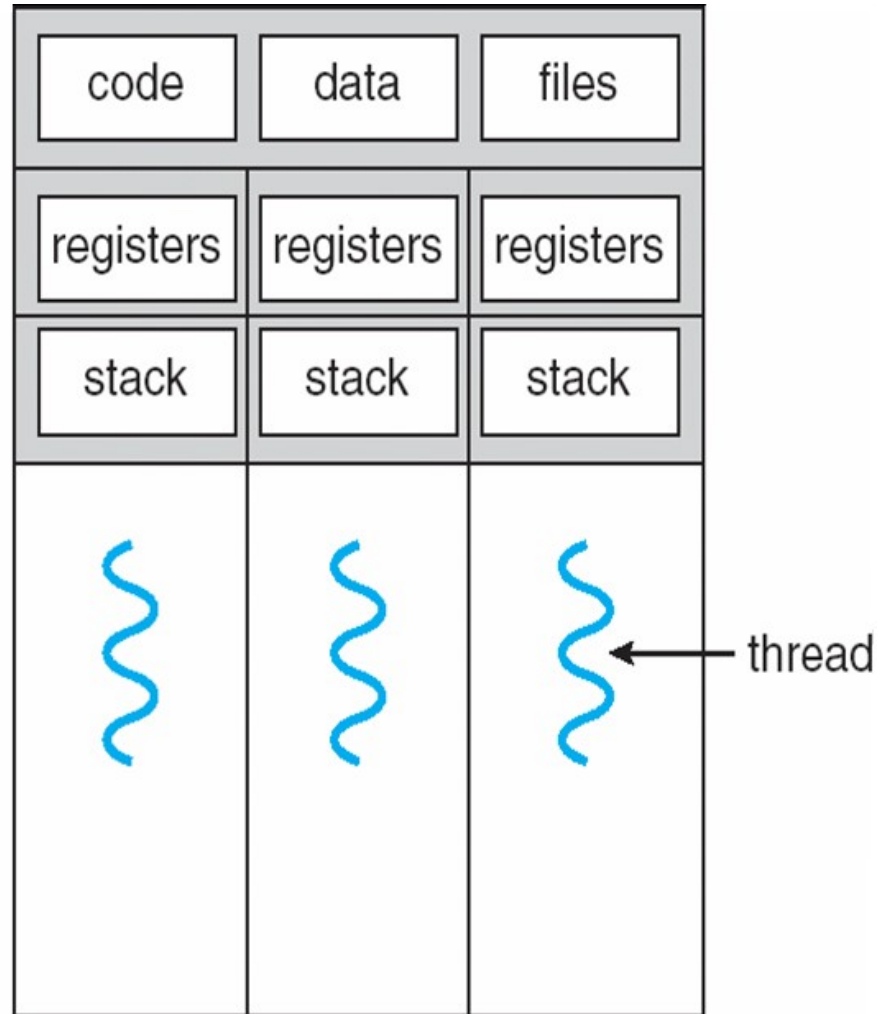
Thread Definition

- A thread is a basic unit of CPU utilization within a process
- Each thread has its own
 - thread ID
 - program counter
 - register set
 - Stack
- It shares the following with other threads within the same process
 - code section
 - data section
 - the heap (dynamically allocated memory)
 - open files and signals
- **Concurrency:** A multi-threaded process can do multiple things at once

The Typical Figure

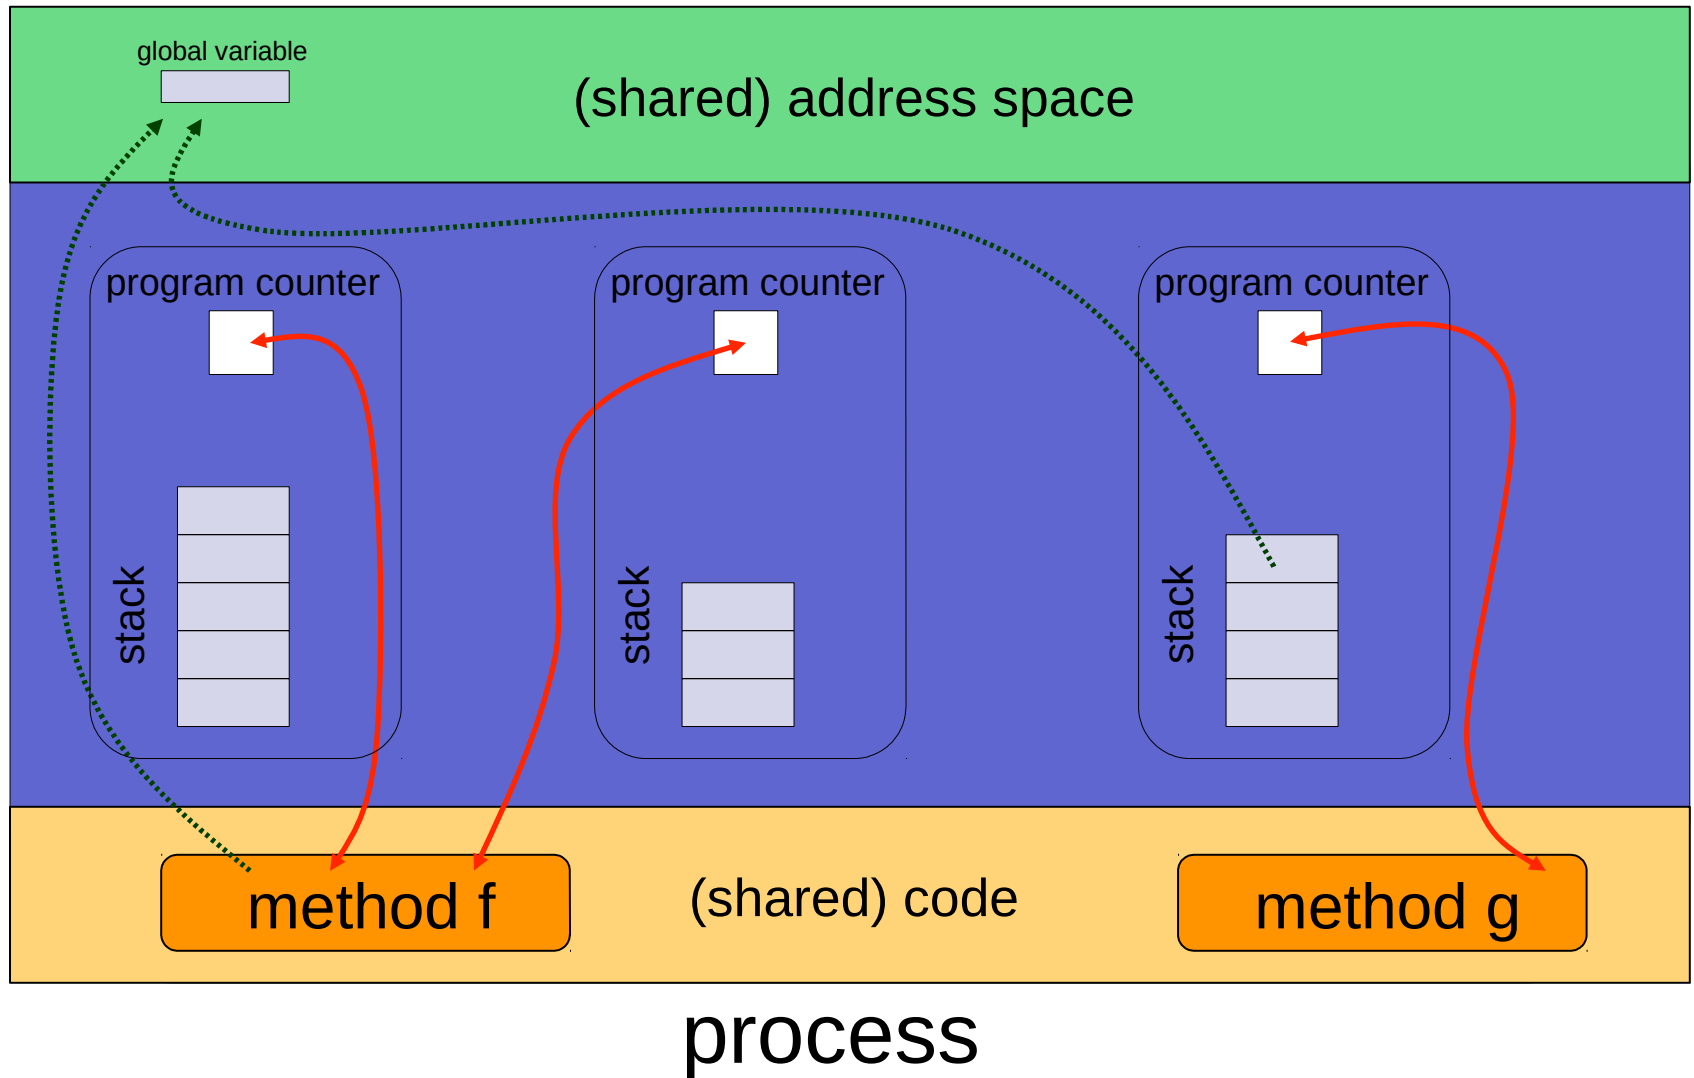


single-threaded process

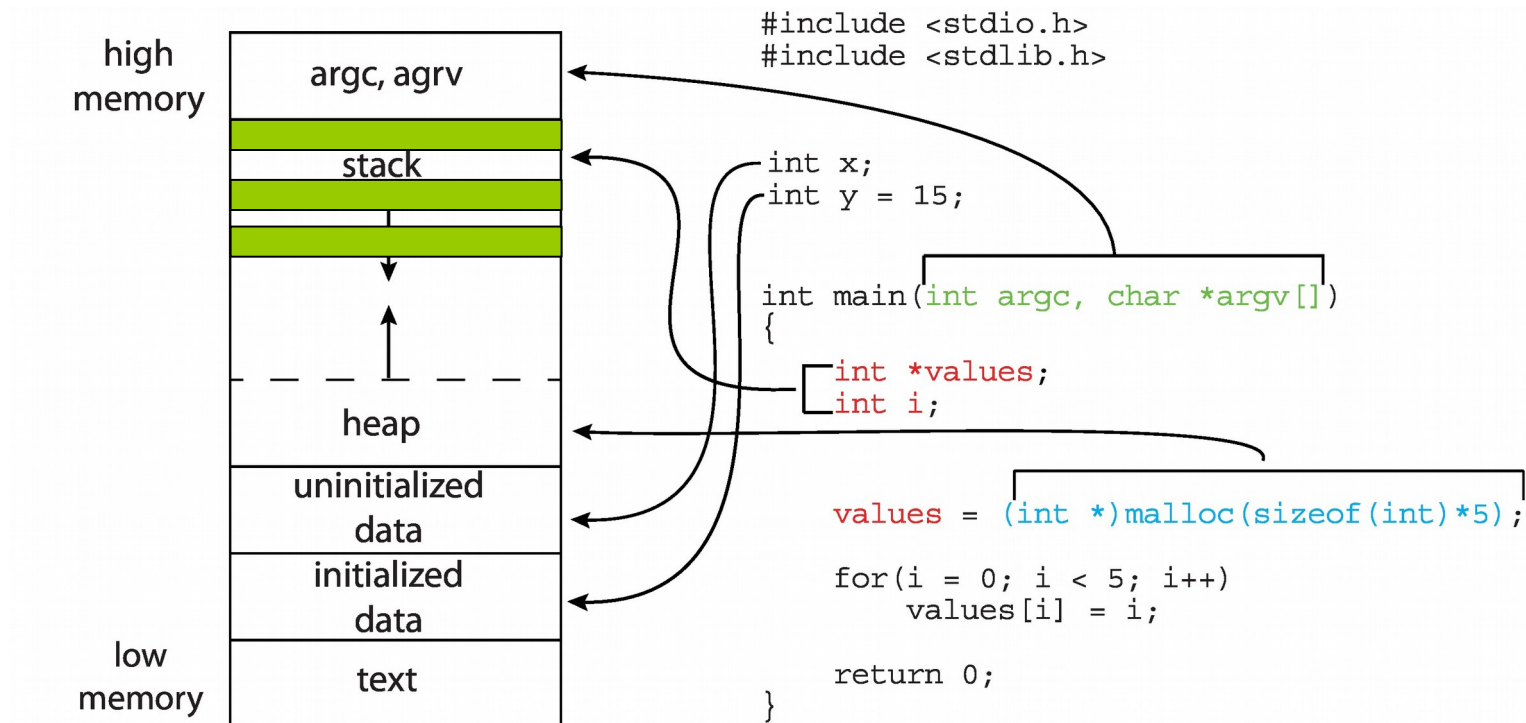


multithreaded process

A More Detailed Figure



Revisit - Memory Layout of a C Program



Multi-Threaded Program

- Source-code view

- a blue thread
- a red thread
- a green thread

[illegible]

```

1  #
2  # Terminal - vim - 101x71
3  #
4  # print_queue()
5  #
6  # print_queue(sb, diffout q) {
7  #   sb->fd->task->id;
8  #   job_descriptor->id;
9  #
10 #   fprintf(stderr, "==== %s\n", task->id);
11 #   sb->fd->foreach(job_descriptor->job_descriptor->id) {
12 #     fprintf(stderr, "%s %s %s\n", descriptor->id, descriptor->nodes,
13 #       descriptor->requested_duration, descriptor->actual_duration);
14 #   }
15 #   printf(stderr, "\n");
16 # }
17 #
18 # start_job()
19 #
20 # void start_job(job_descriptor_t jd, scheduler_bookkeeping_t bk)
21 #
22 # {
23 #   int i;
24 #   time = H5O_get_clock();
25 #
26 #   Note: the job/simulation of a new job */
27 #
28 #   task =
29 #   H5O_task_create(job_scheduler->id, (void*)jd);
30 #   if (i <= task->task->task->id) {
31 #     fprintf(stderr, "Error while sending a job start notification to the job simulator");
32 #   }
33 #   sb->fd->task->id;
34 #
35 #   Note: the simulator */
36 #
37 #   m_task->id;
38 #   int i = jd->id;
39 #   task =
40 #   H5O_task_create(job_scheduler->id, (void*)jd);
41 #   if (i <= task->task->task->id) {
42 #     fprintf(stderr, "Error while sending a job channel started");
43 #   }
44 #   sb->fd->task->id;
45 #
46 #   Note: send simulator of my queue size */
47 #   (the channel is already done anywhere) */
48 #
49 #   m_task->id;
50 #   task =
51 #   H5O_task_create(job_scheduler->id, (void*)queue_size);
52 #   if (i <= task->task->task->id) {
53 #     fprintf(stderr, "Error while sending a job channel queue size");
54 #   }
55 #   sb->fd->task->id;
56 #
57 #   return;
58 # }
59 #
60 # scheduler_init()
61 #
62 # void scheduler_init(job_scheduler_algorithm_t alg, scheduler_bookkeeping_t bk)
63 #
64 # {
65 #   // initialize the job descriptor queues and the number of free nodes */
66 #   bk->queue = sb->fd->queue;
67 #   bk->running = sb->fd->running;
68 #
69 #   // Grouping for initialization */
70 #
71 #   case FBFS:
72 #     scheduler_init_fbfs(bk);
73 #   }
74 # }

```

```

        }
        break;
    }

    if ((num_target_schedulers > 1)
        pending_job->flooded_across == 1)
    {
        free(pruned_schedulers);
        return;
    }

    #
    # find_Rn_target_schedulers()
    #
    void find_Rn_target_schedulers(int x, scheduler_info_t preferred,
                                   scheduler_info_t *schedulers, int num_schedulers,
                                   int *num_target_schedulers, scheduler_info_t **target_schedulers)
    {
        int i;
        int n;
        done;

        num_target_schedulers = 0;

        for (i = 0; i < (MIN(x, num_schedulers)); i++) {
            if (i == 0) { // first pass: preferred scheduler +/
                num_schedulers++;
                schedulers[i] = preferred;
                n = 1;
                break;
            }
            else
            {
                while (1) {
                    n = random_integer_biased(1, num_schedulers - 1);
                    n = random_integer(0, num_schedulers - 1);

                    done = 1;
                    for (j = 0; j < i; j++)
                        if ((target_schedulers[j]) == schedulers[n])
                            done = 0;
                    if (done)
                        break;
                }

                num_target_schedulers++;
                *target_schedulers = RBLLOC(target_schedulers,
                                             (num_target_schedulers) * sizeof(scheduler_info_t));
                target_schedulers[num_target_schedulers - 1] = schedulers[n];
            }
        }

        return;
    }

    #
    # find_Sx_target_schedulers()
    #
    void find_Sx_target_schedulers(int x, scheduler_info_t preferred,
                                   scheduler_info_t *schedulers, int num_schedulers,
                                   int *num_target_schedulers, scheduler_info_t **target_schedulers)
    {
        scheduler_info_t *tosort;
        int i;
    }

```

```

Terminal — vim — 96x42
#include "receiver.h"
#include "stdlib.h"

XBT_LOG_NEW_DEFAULT_CATEGORY(RECEIVER, "Logging for the receiver process");

/* Receiver function
 * arg #1: port
 * arg #2: dynar name
 */
int receiver(int argc, char *argv[])
{
    int port = 1;
    xbt_dict_t dict;
    dynar_t dynar;

    /* Parse the first argument */
    if (argc != 1, *argv[0]) {
        xbt_assert(0, "Invalid port '%s' for a receiver process", argv[0]);
    }

    /* Get the dynar */
    if ((dynar = (dynar_t)xbt_dict_get_or_null(receiver_dynar_dict, (const char *)argv[0])) != 0) {
        xbt_assert(0, "Cannot find dynar '%s' for a receiver process", argv[0]);
    }

    /* Loop */
    while (1) {
        int i;
        e_t task = NULL;

        if (MSG_task_get(task, port) != MSG_OK) {
            xbt_assert(0, "Error while receiving a task in a receiver process");
        }

        /* Pick the task at a random location in the dynar */
        i = rand() % (dynar ? dynar->length : 0);
        dynar_t task_dynar = (dynar_t)xbt_dict_get(dynar, (void *)task);
    }
}

```


Advantages of Threads?

■ Economy:

- Creating a thread is cheap
 - ▶ Much cheaper than creating a process
- Context-switching between threads is cheap
 - ▶ Much cheaper than between processes

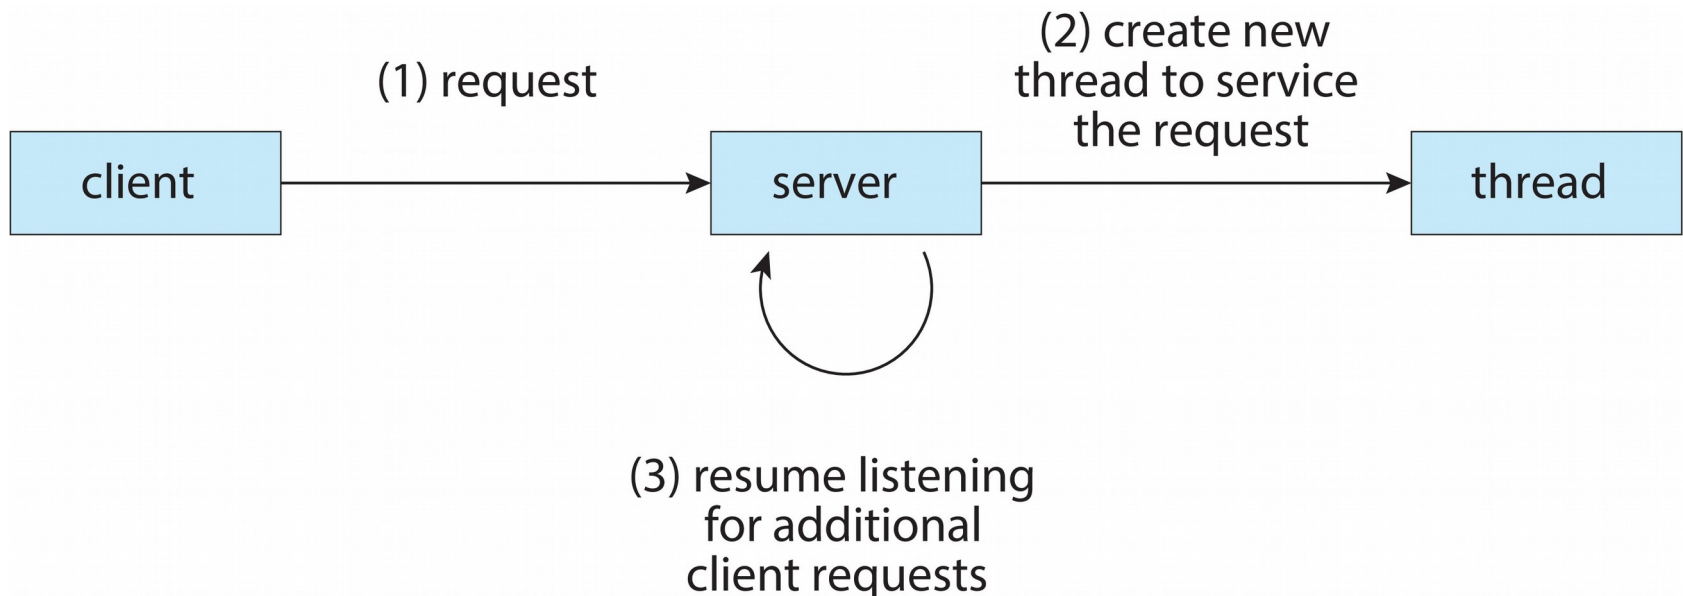
■ Resource Sharing:

- Threads naturally share memory
 - ▶ With processes you have to use possibly complicated IPC (e.g., Shared Memory Segments)
- Having concurrent activities in the same address space is very powerful
 - ▶ But fraught with danger

Advantages of Threads?

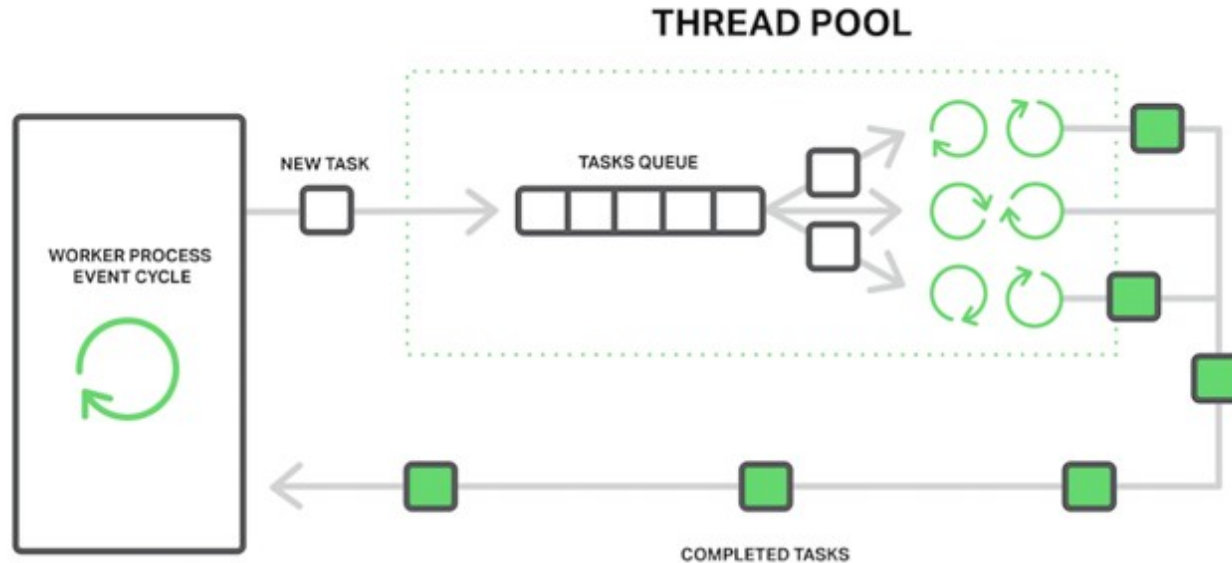
■ Responsiveness

- A program that has concurrent activities is more responsive
 - ▶ While one thread blocks waiting for some event, another can do something
 - ▶ e.g. Spawn a thread to answer a client request in a client-server implementation
- This is true of processes as well, but with threads we have better sharing and economy



Advantages of Threads?

- Thread Pools in NGINX Boost Performance 9x
 - nginx : worker process -> thread pool



Advantages of Threads?

■ Responsiveness

- A program that has concurrent activities is more responsive
 - ▶ While one thread blocks waiting for some event, another can do something
 - ▶ e.g. Spawn a thread to answer a client request in a client-server implementation
- This is true of processes as well, but with threads we have better sharing and economy

■ Scalability

- Running multiple “threads” at once uses the machine more effectively
 - ▶ e.g., **on a multi-core machine**
- This is true of processes as well, but with threads we have better sharing and economy

Drawbacks of Threads

- One drawback of thread-based concurrency compared to process-based concurrency: If one thread fails (e.g., a segfault), then the process fails
 - And therefore the whole program
- This leads to process-based concurrency
 - e.g., The Google Chrome Web browser
 - See <http://www.google.com/googlebooks/chrome/>
- Sort of a throwback to the pre-thread era
 - Threads have been available for 20+ years
 - Very trendy recently due to multi-core architectures

Drawbacks of Threads

- Threads may be more memory-constrained than processes
 - Due to OS limitation of the address space size of a single process
 - Not a problem any more on 64-bit architecture
- Threads do not benefit from memory protection
 - Concurrent programming with Threads is hard
 - ▶ But so is it with Processes and Shared Memory Segments

Threads on My Machine?

- Let's run `ps aux` and look at several applications
 - `ps aux` and `ps -T -p PID`
 - Chrome
 - Terminal
 - ...

Multi-Threading Challenges


- Typical challenges of multi-threaded programming
 - Dividing activities among threads
 - Balancing load among threads
 - Split data among threads
 - Deal with data dependency and synchronization
 - Testing and Debugging

User Threads vs. Kernel Threads


- Threads can be supported solely in User Space
 - Threads are managed by some user-level thread library (e.g., Java Green Threads)
- Threads can also be supported in Kernel Space
 - The kernel has data structure and functionality to deal with threads
 - Most modern OSes support kernel threads
 - ▶ In fact, Linux doesn't really make a difference between processes and threads (same data structure)

Many-to-One Model

- Advantage: multi-threading is efficient and low-overhead
 - No syscalls to the kernel
- **Major Drawback #1:** cannot take advantage of a multi-core architecture!
- **Major Drawback #2:** if one threads blocks, then all the others do!
- Examples (User-level Threads):
 - Java Green Threads
 - GNU Portable Threads




One-to-One Model




- Removes both drawbacks of the Many-to-One Model
- Creating a new threads requires work by the kernel
 - Not as fast as in the Many-to-One Model
- Example:
 - Linux
 - Windows
 - Solaris 9 and later

Many-to-Many Model

- A compromise
- If a user thread blocks, the kernel can create a new kernel threads to avoid blocking all user threads
- A new user thread doesn't necessarily require the creation of a new kernel thread
- True concurrency can be achieved on a multi-core machine
- Examples:
 - Solaris 9 and earlier
 - Win NT/2000 with the ThreadFiber package



Two-Level Model



- The user can say: “Bind this thread to its own kernel thread”
- Example:
 - IRIX, HP-UX, Tru64 UNIX
 - Solaris 8 and earlier

Thread Libraries

- Thread libraries provide users with ways to create threads in their own programs
 - In C/C++: Pthreads
 - ▶ Implemented by the kernel
 - In C/C++: OpenMP
 - ▶ A layer above Pthreads for convenient multithreading in “easy” cases
 - In Java: Java Threads
 - ▶ Implemented by the JVM, which relies on threads implemented by the kernel

Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)

Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```


Pthreads Example

```
/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
```

Windows Multithreaded C Program

```
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```

OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel

#pragma omp parallel

Create as many threads as there are cores

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}

#pragma omp parallel for
for (i = 0; i < N; i++) {
    c[i] = a[i] + b[i];
}
```

Java Threads

- All memory-management headaches go away with Java Threads
 - In nice Java fashion
- Several programming languages have long provided constructs/abstractions for writing concurrent programs
 - Modula, Ada, etc.
- Java threads may be created by:

- Extending Thread class


```
class MyThread extends Thread {  
    public void run() {  
        . . .  
    }  
}  
MyThread t = new MyThread();
```

- Implementing the Runnable interface

```
public interface Runnable  
{  
    public abstract void run();  
}
```

Thread Scheduling

- The JVM keeps track of threads, enacts the thread state transition diagram
- Question: who decides which runnable thread to run?
- Old versions of the JVM used **Green Threads**
 - User-level threads implemented by the JVM
 - Invisible to the O/S




Beyond Green Threads

- Green threads have all the disadvantages of user-level threads (see earlier)
 - Most importantly: Cannot exploit multi-core, multi-processor architectures
- The JVM now provides **native threads**
 - Green threads are typically not available anymore
 - you can try to use “java -green” and see what your system says

Java Threads / Kernel Threads

- In modern JVMs, application threads are *mapped* to kernel threads



Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations

Semantics of `fork()` and `exec()`

- What happens when a thread calls `fork()`?
- Two possibilities:
 - A new process is created that has only one thread (the copy of the thread that called `fork()`), or
 - A new process is created with all threads of the original process (a copy of all the threads, including the one that called `fork()`)
- Some OSes provide both options
 - In Linux the first option above is used
- If one calls `exec()` after `fork()`, all threads are “wiped out” anyway

Signals

- We've talked about signals for processes
 - Signal handlers are either default or user-specified
 - `signal()` and `kill()` are the system calls
- In a multi-threaded program, what happens?
- Multiple options
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals

Signals

- Most UNIX versions: a thread can say which signals it accepts and which signals it doesn't accept
- On Linux: dealing with threads and signals is tricky but well understood with many tutorials on the matter and man pages
 - `man pthread_sigmask`
 - `man sigemptyset`
 - `man sigaction`

Safe Thread Cancellation

- One potentially useful feature would be for a thread to simply terminate another thread
- Two possible approaches:
 - **Asynchronous** cancellation
 - ▶ One thread terminates another immediately
 - **Deferred** cancellation
 - ▶ A thread periodically checks whether it should terminate

Safe Thread Cancellation

■ Two possible approaches:

- **Asynchronous** cancellation
 - ▶ One thread terminates another immediately
- **Deferred** cancellation
 - ▶ A thread periodically checks whether it should terminate

■ Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```

Safe Thread Cancellation

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches **cancellation point**
 - ▶ I.e. **pthread_testcancel()**
 - ▶ Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals

Safe Thread Cancellation

- The problem with asynchronous cancellation:
 - may lead to an inconsistent state or to a synchronization problem if the thread was in the middle of “something important”
 - Absolutely terrible bugs lurking in the shadows
- The problem with deferred cancellation: the code is cumbersome due to multiple cancellation points
 - should I die? should I die? should I die?
- In Java, the `Thread.stop()` method is deprecated, and so cancellation has to be deferred

Operating System Examples

- Windows Threads
- Linux Threads


Windows Threads

- Windows API – primary API for Windows applications
- Implements the one-to-one mapping, kernel-level
- Each thread contains
 - A thread id
 - Register set representing state of processor
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- The register set, stacks, and private storage area are known as the **context** of the thread

Windows Threads (Cont.)

- The primary data structures of a thread include:
 - ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space
 - KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
 - TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space

Windows Threads Data Structures



Linux Threads

- Linux does not distinguish between processes and threads: they're called **tasks**
- Kernel data structure: task_struct

```
591
592 struct task_struct {
593 #ifdef CONFIG_THREAD_INFO_IN_TASK
594     /*
595      * For reasons of header soup (see current_thread_info()), this
596      * must be the first element of task_struct.
597      */
598     struct thread_info          thread_info;
599 #endif
600     /* -1 unrunnable, 0 runnable, >0 stopped: */
601     volatile long               state;
602
603     /*
604      * This begins the randomizable portion of task_struct. Only
605      * scheduling-critical items should be added above here.
606      */
607     randomized_struct_fields_start
608
609     void                        *stack;
610     atomic_t                    usage;
611     /* Per task flags (PF_*), defined further below: */
612     unsigned int                flags;
613     unsigned int                ptrace;
614
```

Linux Threads

- In Linux, a thread is also called a light-weight process (LWP)
- The `clone()` syscall is used to create a task
 - Shares execution context with its parent
 - pthread library uses `clone()` to implement threads. Refer to `./nptl/sysdeps/pthread/createthread.c`


flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.

Linux Threads


■ Single-threaded process vs multi-threaded process

```
wenbo@wenbo-desktop:~/KERNEL/linux.git$ ps -eLf
```

UID	PID	PPID	LWP	C	NLWP	STIME	TTY	TIME	CMD
root	1	0	1	0	1	3月11 ?		00:00:19	/sbin/init splash
root	2	0	2	0	1	3月11 ?		00:00:00	[kthreadd]
root	4	2	4	0	1	3月11 ?		00:00:00	[kworker/0:0H]
root	6	2	6	0	1	3月11 ?		00:00:00	[mm_percpu_wq]
root	7	2	7	0	1	3月11 ?		00:00:00	[ksoftirqd/0]
root	8	2	8	0	1	3月11 ?		00:00:31	[rcu_sched]
root	9	2	9	0	1	3月11 ?		00:00:00	[rcu_bh]
root	10	2	10	0	1	3月11 ?		00:00:00	[migration/0]
root	11	2	11	0	1	3月11 ?		00:00:00	[watchdog/0]
root	704	1	704	0	1	3月11 ?		00:00:00	/usr/sbin/cron -f
root	718	1	718	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	882	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	883	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	884	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	885	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	917	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	921	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	922	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	923	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	924	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd



single-threaded process



multithreaded process

Linux Threads

■ Single-threaded process vs multi-threaded process

```
wenbo@wenbo-desktop:~/KERNEL/linux.git$ ps -eLf
```


UID	PID	PPID	LWP	C	NLWP	STIME	TTY	TIME	CMD
root	1	0	1	0	1	3月11 ?		00:00:19	/sbin/init splash
root	2	0	2	0	1	3月11 ?		00:00:00	[kthreadd]
root	4	2	4	0	1	3月11 ?		00:00:00	[kworker/0:0H]
root	6	2	6	0	1	3月11 ?		00:00:00	[mm_percpu_wq]
root	7	2	7	0	1	3月11 ?		00:00:00	[ksoftirqd/0]
root	8	2	8	0	1	3月11 ?		00:00:31	[rcu_sched]
root	9	2	9	0	1	3月11 ?		00:00:00	[rcu_bh]
root	10	2	10	0	1	3月11 ?		00:00:00	[migration/0]
root	11	2	11	0	1	3月11 ?		00:00:00	[watchdog/0]
root	704	1	704	0	1	3月11 ?		00:00:00	/usr/sbin/cron -f
root	718	1	718	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	882	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	883	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	884	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	885	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	917	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	921	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	922	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	923	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	924	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd

```
787  /* PID/PID hash table linkage. */
788  struct pid                *thread_pid;
789  struct hlist_node         pid_links[PIDTYP
790  struct list_head          thread_group;
791  struct list_head          thread_node;
792
793  struct completion         *vfork_done;
794
795  /* CLONE_CHILD_SETTID: */
796  int __user                *set_child_tid;
---
```

Linux Threads

- Linux does not distinguish between processes and threads: they're called **tasks**
 - Kernel data structure: `task_struct`
- A process is
 - either a single thread + an address space
 - ▶ PID is thread ID
 - or multiple threads + an address space
 - ▶ PID is the leading thread ID

Threads within Process



Threads with Process – What is shared

PROCESS	THREAD	TSK=ffff8c4c4bf3c5c0	PID=718	STACK=ffff985c82268000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d52e80	PID=882	STACK=ffff985c82390000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d545c0	PID=883	STACK=ffff985c822e8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c491b45c0	PID=884	STACK=ffff985c8218c000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4cbeb1740	PID=885	STACK=ffff985c821ec000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4ae1ae80	PID=917	STACK=ffff985c823c8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b562e80	PID=921	STACK=ffff985c82418000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c48340000	PID=922	STACK=ffff985c823b0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c472bae80	PID=923	STACK=ffff985c821f4000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b5945c0	PID=924	STACK=ffff985c81fa8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46775d00	PID=925	STACK=ffff985c822a8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b692e80	PID=973	STACK=ffff985c82438000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b78ae80	PID=974	STACK=ffff985c823c0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46e1dd00	PID=975	STACK=ffff985c824b8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840

Threads within Process – What is shared

PROCESS	THREAD	TSK=ffff8c4c4bf3c5c0	PID=718	STACK=ffff985c82268000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d52e80	PID=882	STACK=ffff985c82390000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d545c0	PID=883	STACK=ffff985c822e8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c491b45c0	PID=884	STACK=ffff985c8218c000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4beb1740	PID=885	STACK=ffff985c821ec000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4ae1ae80	PID=917	STACK=ffff985c823c8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b562e80	PID=921	STACK=ffff985c82418000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c48340000	PID=922	STACK=ffff985c823b0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c472bae80	PID=923	STACK=ffff985c821f4000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b5945c0	PID=924	STACK=ffff985c81fa8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46775d00	PID=925	STACK=ffff985c822a8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b692e80	PID=973	STACK=ffff985c82438000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b78ae80	PID=974	STACK=ffff985c823c0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46e1dd00	PID=975	STACK=ffff985c824b8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
		task_struct	pid	stack	comm	mm_struct	

Not Shared

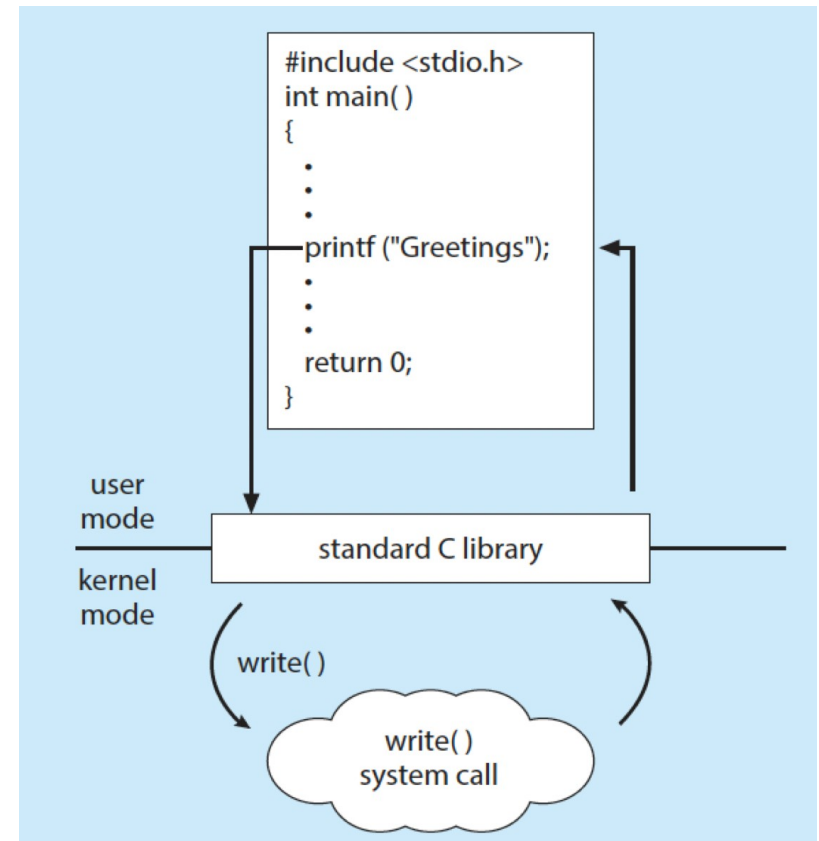
Shared

User thread to kernel thread mapping

- One task
 - One task struct – PCB
 - Can be executed in user space
 - ▶ User code, user space stack
 - Can be executed in kernel space
 - ▶ Such as calls a system call
 - ▶ Execution flow traps to kernel
 - ▶ Execute kernel code, use kernel space stack

User thread to kernel thread mapping

- One task
 - One task struct – PCB
 - Execute in one thread
 - Can be executed in user space
 - ▶ User code, user space stack
 - Can be executed in kernel space
 - ▶ Such as calls a system call
 - ▶ Execution flow traps to kernel
 - ▶ Execute kernel code, use kernel space stack

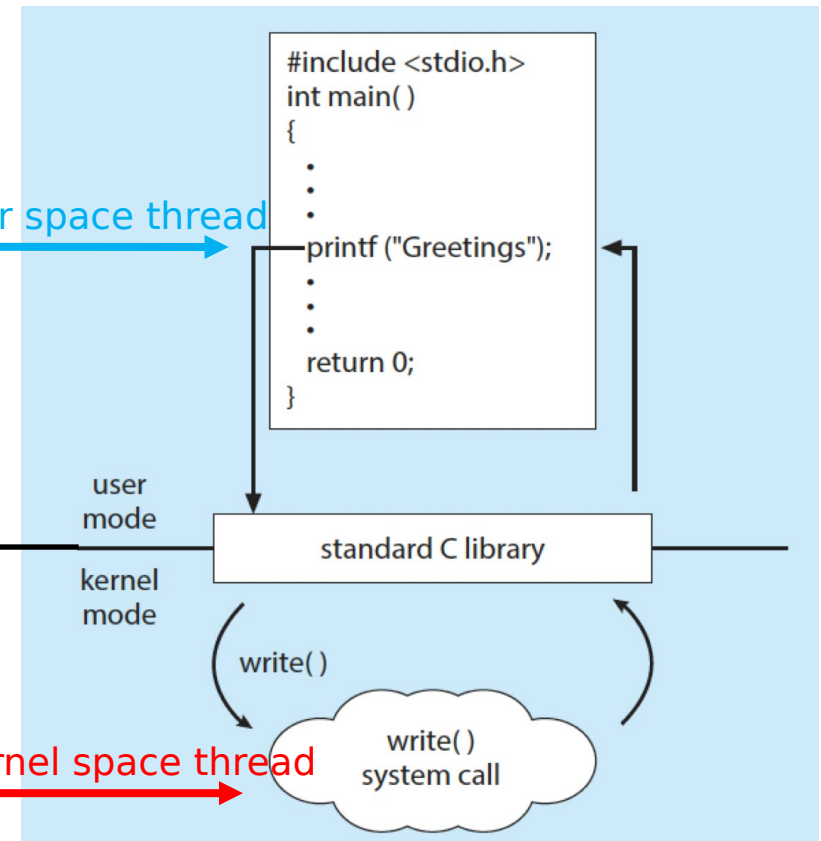


User thread to kernel thread mapping

■ One task

- One task struct – PCB
- Execute in one thread
- Can be executed in user space
 - ▶ User code, user space stack

User space thread



- Can be executed in kernel space
 - ▶ Such as calls a system call
 - ▶ Execution flow traps to kernel
 - ▶ Execute kernel code, use kernel space stack

User thread to kernel thread mapping

- One task in Linux
 - Same task_struct (PCB) means same thread
 - ▶ Also viewed as 1:1 mapping
 - ▶ One user thread maps to one kernel thread
 - ▶ But actually, they are the same thread
 - Can be executed in user space
 - ▶ User code, user space stack
 - Can be executed in kernel space
 - ▶ Kernel code, kernel space stack
- Kernel thread also uses to represent threads that has no user space part
 - Such as kernel wants to flush dirty buffer to disk
 - It creates a thread, running only in kernel mode

Takeaway

- Thread is the basic execution unit
 - Has its own registers, pc, stack
- Thread vs Process
 - What is shared and what is not
- Pros and cons of thread

- Lab1 of both tracks are out