


# Virtual Memory and Linux


---

Based on Alan Ott's slides on  
Embedded Linux Conference  
April 4-6, 2016



---

# Physical Memory



# Flat Memory

---

- Older and modern, but simple systems have a single address space
  - Memory and peripherals share
    - Memory will be mapped to one part
    - Peripherals will be mapped to another
  - All processes and OS share the same memory space
    - No memory protection!
    - User space can stomp kernel mem!

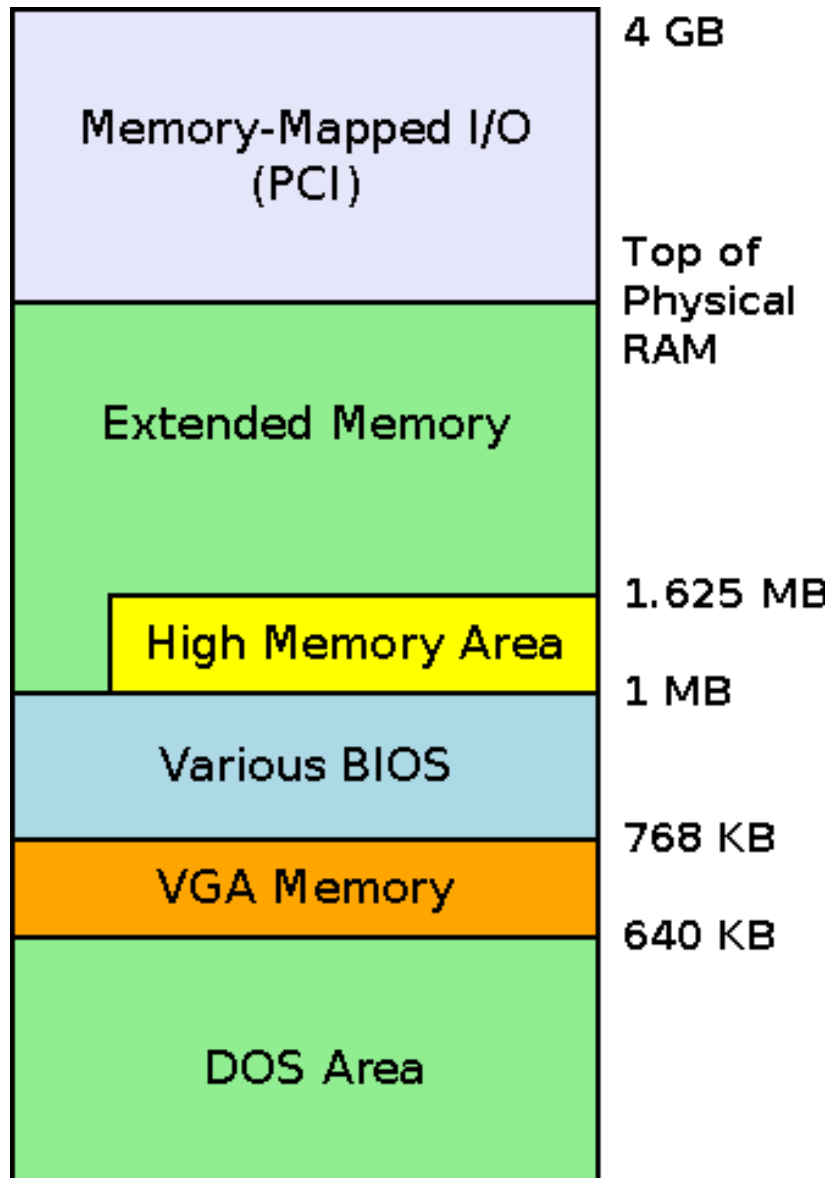
# Flat Memory

---

- CPUs with flat memory
  - 8086-80206
  - ARM Cortex-M
  - IoT chips - most 8- and 16-bit systems

# x86 Memory Map

---



- Lots of Legacy
- RAM is split (DOS Area and Extended)
- Hardware mapped between RAM areas.
- High and Extended accessed differently

# Limitations

---

- Portable C programs expect flat memory
  - Accessing memory by segments limits portability
- Management is tricky
  - Need to keep processes separated
- No protection
  - Rogue programs can corrupt the entire system

---

# Virtual Memory



# What is Virtual Memory?

---

- Virtual Memory is an address mapping
  - Maps virtual address space to physical address space
    - Maps virtual addresses to physical RAM
    - Maps virtual addresses to hardware devices
      - PCI devices
      - GPU RAM
      - On-SoC IP blocks



# What is Virtual Memory?

---

- Advantages
  - Each processes can have a different memory mapping
    - One process's RAM is inaccessible (and invisible) to other processes.
      - Built-in memory protection
    - Kernel RAM is invisible to userspace processes
  - Memory can be moved
  - Memory can be swapped to disk

# What is Virtual Memory?

---

- Advantages (cont)
  - Hardware device memory can be mapped into a process's address space
    - Requires kernel perform the mapping
  - Physical RAM can be mapped into multiple processes at once
    - Shared memory, shared libraries
  - Memory regions can have access permissions
    - Read, write, execute

# Virtual Memory Details

---

- Two address spaces
  - Physical addresses
    - Addresses as used by the hardware (except CPU)
      - DMA, peripherals
  - Virtual addresses
    - Addresses as used by software (CPU generated)
      - Load/Store instructions (RISC)
      - Any instruction accessing RAM (CISC)

# Virtual Memory Details

---

- Mapping is performed in hardware
  - No performance penalty for accessing already-mapped RAM regions
  - Permissions are handled without penalty
  - The same CPU instructions are used for accessing RAM and mapped hardware
  - Software, during its normal operation, will only use virtual addresses.
    - Includes kernel and user space

# Memory-Management Unit

---

- The memory-management unit (MMU) is the hardware responsible for implementing virtual memory.
  - Sits between the CPU core and memory
  - Most often part of the physical CPU itself.
    - On ARM, it's part of the licensed core.
  - Separate from the RAM controller
    - DDR controller is a separate IP block

# Memory-Management Unit

---

- MMU (cont)
  - Transparently handles all memory accesses from Load/Store instructions
    - Maps accesses using virtual addresses to system RAM
    - Maps accesses using virtual addresses to memory-mapped peripheral hardware
    - Handles permissions
    - Generates an exception (page fault) on an invalid access
      - Unmapped address or insufficient permissions

# Translation Lookaside Buffer

---

- The TLB stores the mappings from virtual to physical address space in hardware
  - Also holds permission bits
- TLB is a part of the MMU
- TLB is consulted by the MMU when the CPU accesses a virtual address

# Page Faults

---

- A page fault is a CPU exception, generated when software attempts to use an invalid virtual address. There are three cases:
  - The virtual address is not mapped for the process requesting it.
  - The processes has insufficient permissions for the address requested.
  - The virtual address is valid, but swapped out
    - This is a software condition

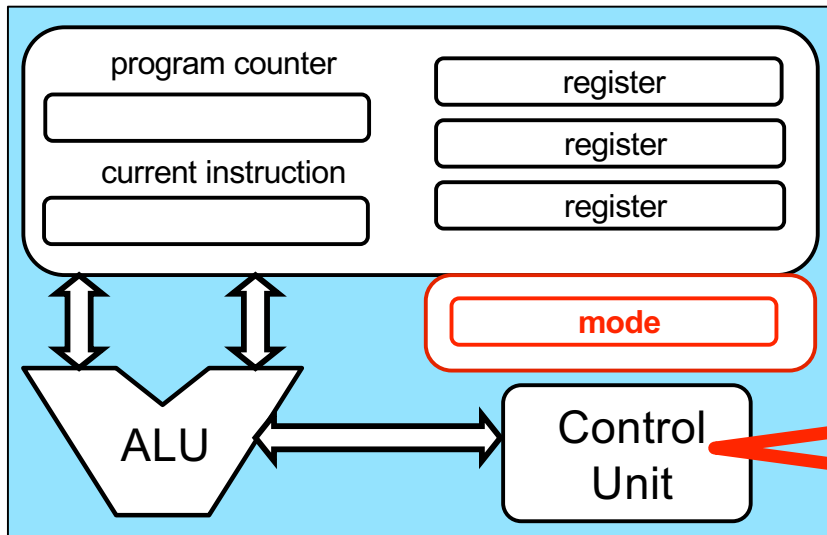


# Lazy Allocation

---

- The kernel uses lazy allocation of physical memory.
  - When memory is requested by userspace, physical memory is not allocated until it's touched.
  - This is an optimization, knowing that many userspace programs allocate more RAM than they ever touch.
    - Buffers, etc.

# Kernel address space vs user address space



- Decode instruction
- Determine if instruction is privileged or not
  - Based on the instruction code (e.g., the binary code for all privileged instructions could start with '00')
- If instruction is privileged and mode == user, then abort!
  - Raise a "trap"

# Kernel address space vs user address space

---

- Address space can be **privileged** or **non-privileged** too
- For address space (AS) isolation, dividing AS into two parts
  - Kernel AS (**privileged**) and user AS (**non-privileged**)
  - Kernel code can access **both** AS
  - User code can **only** access user AS, **cannot** access kernel AS

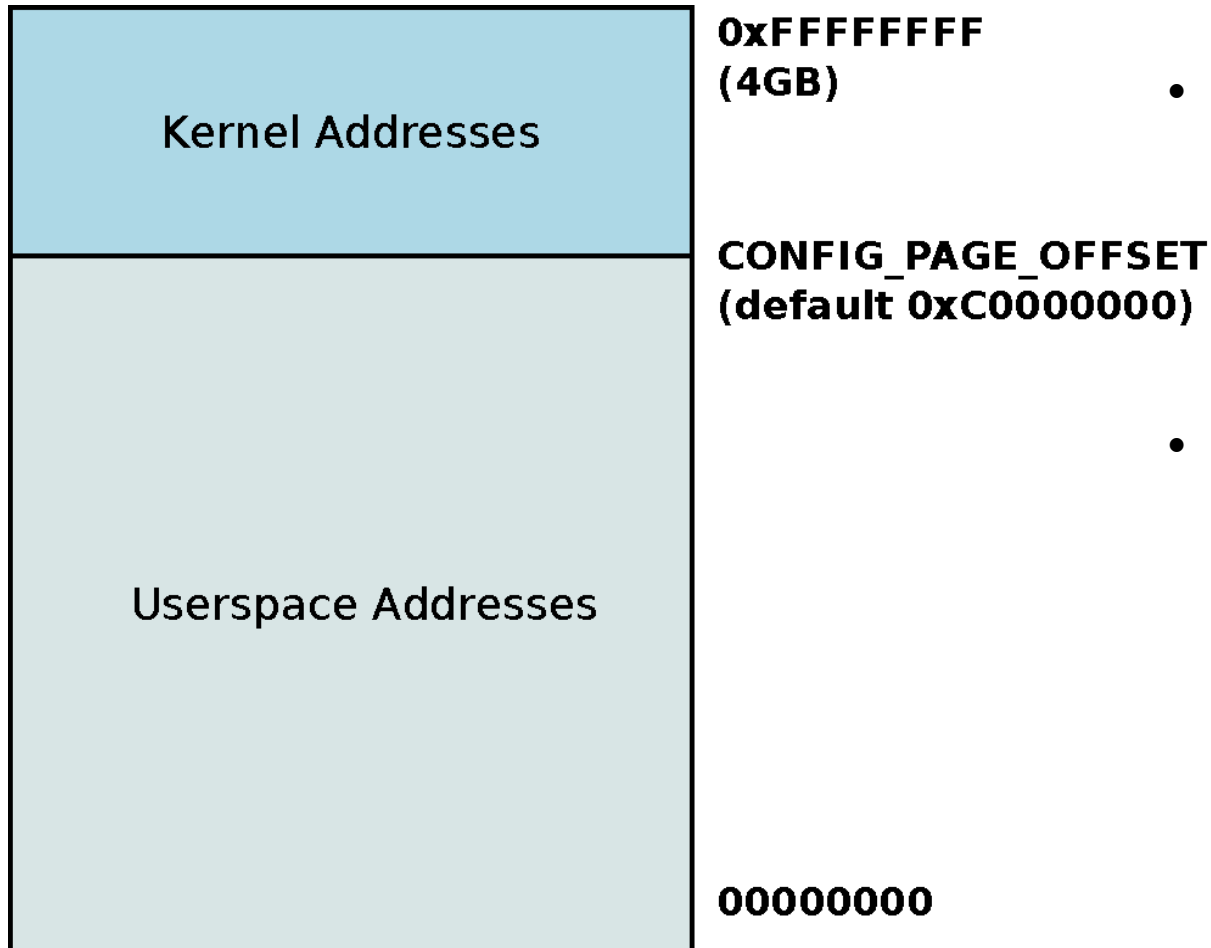
# Virtual Addresses

---

- In Linux, the kernel uses virtual addresses, as userspace processes do.
- Virtual address space is split.
  - The upper part is used for the kernel
  - The lower part is used for userspace
- On 32-bit, the split is at  
0xC0000000

# Virtual Addresses - Linux

---



- By default, the kernel uses the top 1GB of virtual address space.
- Each userspace processes get the lower 3GB of virtual address space.

# Virtual Addresses - Linux

---

- Kernel address space is the area above `CONFIG_PAGE_OFFSET`.
  - For 32-bit, this is configurable at kernel build time.
    - The kernel can be given a different amount of address space as desired
      - See `CONFIG_VMSPLIT_1G`, `CONFIG_VMSPLIT_2G`, etc.
  - For 64-bit, the split varies by architecture, but it's high enough
    - `0x8000000000000000`– ARM
    - `0xffff880000000000`– x86\_64

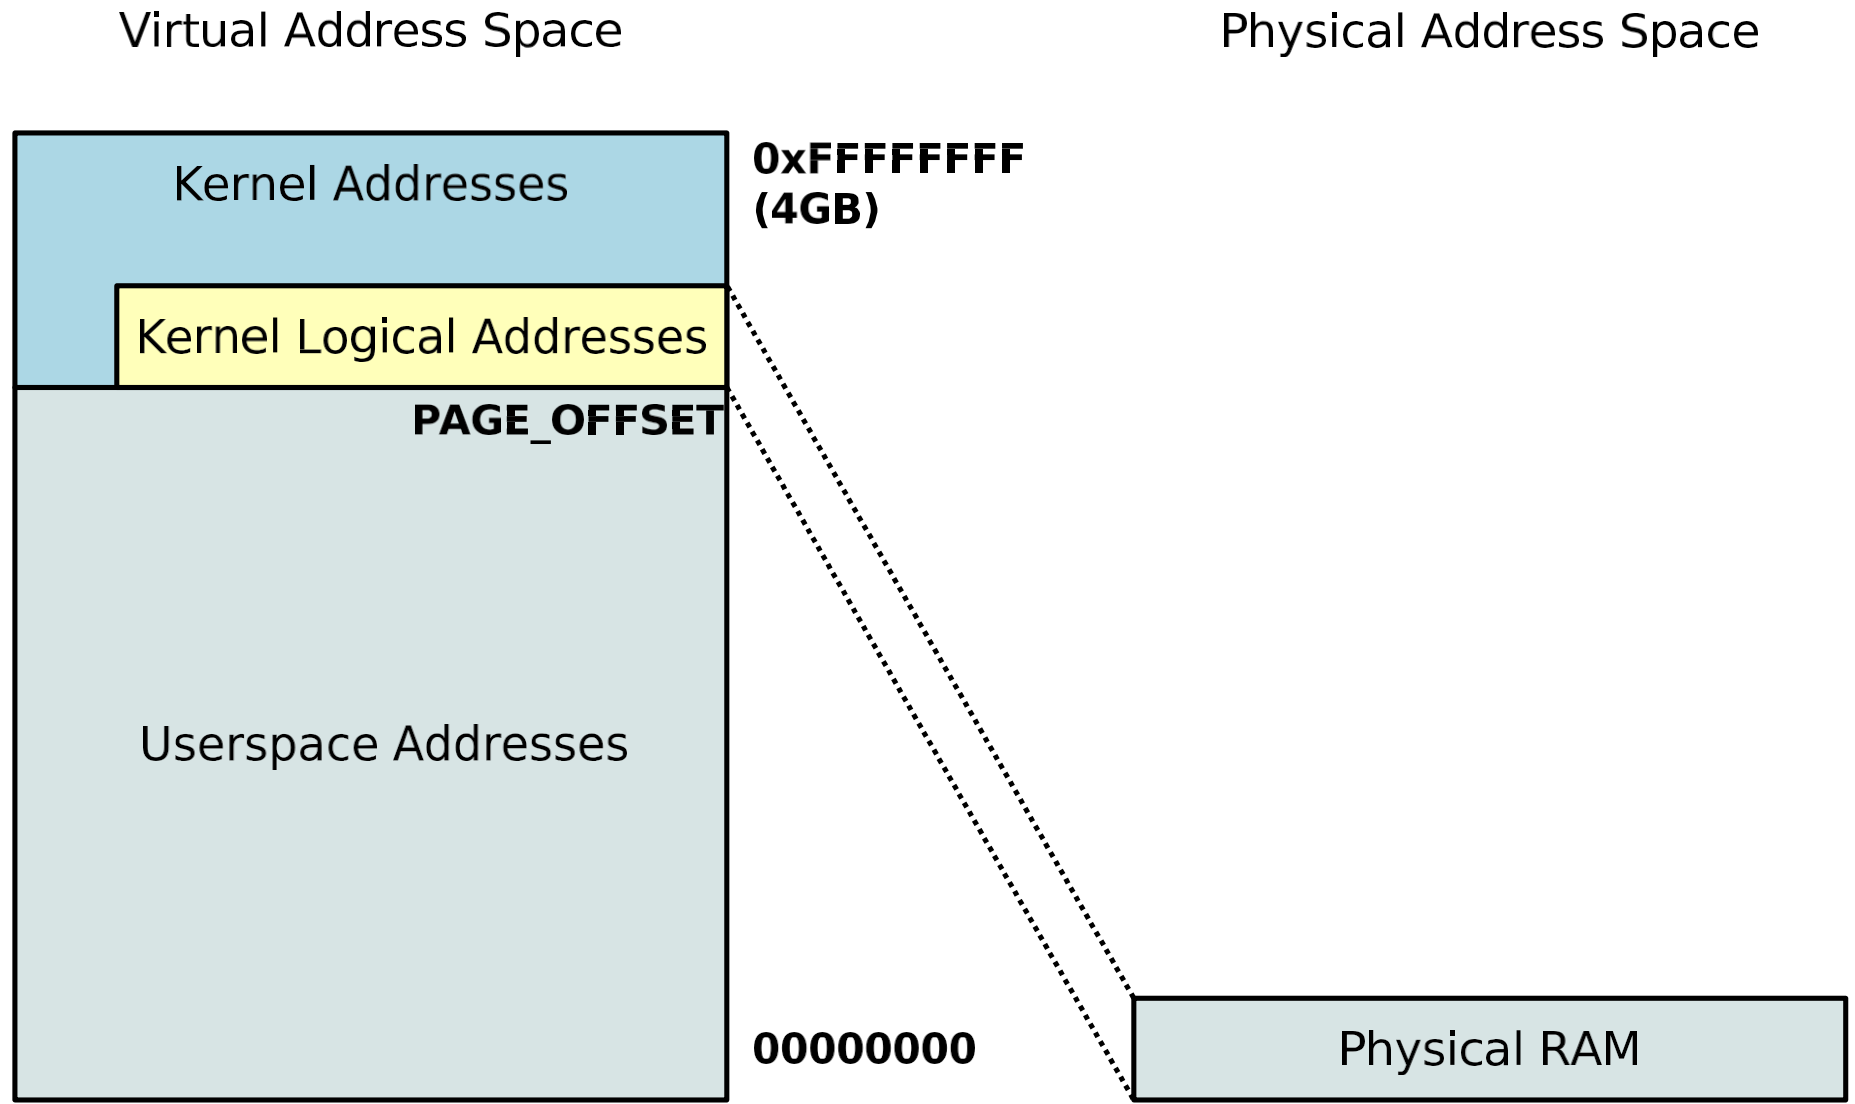
# Virtual address and physical RAM mapping

---

- How kernel manage physical RAM?
  - Such as allocate a physical frame to user process
- Must map physical RAM to kernel AS
  - Must have an address in kernel AS
  - For 32-bit, kernel address space (**AS**) is 1GB
    - No problem to handle small RAM ( $\leq 896\text{MB}$ )

# Kernel Logical Addresses

---






# Kernel Logical Addresses

---

- Kernel Logical Addresses
  - Addresses above PAGE\_OFFSET
  - First 896MB of the kernel AS
  - Virtual addresses are a linear offset from their physical addresses.
    - Eg: Virt: 0xc0000000 → Phys: 0x00000000
  - This makes converting between physical and virtual addresses easy

# How to do VA to PA translation?

- Given a virtual address  $v$ , what is the physical address?
- Given a physical address  $p$ , what is the virtual address?



# Kernel Logical Addresses

---

- Kernel logical address space includes:
  - Memory allocated with `kmalloc()` and most other allocation methods
  - Kernel stacks (per process)
- Kernel logical memory can never be swapped out!

# Kernel Logical Addresses

---

- Kernel Logical Addresses use a linear mapping between physical and virtual address space.
- This means virtually-contiguous regions are by nature also physically contiguous
  - This, combined with the inability to be swapped out, makes them suitable for DMA transfers.

# Large physical RAM

---

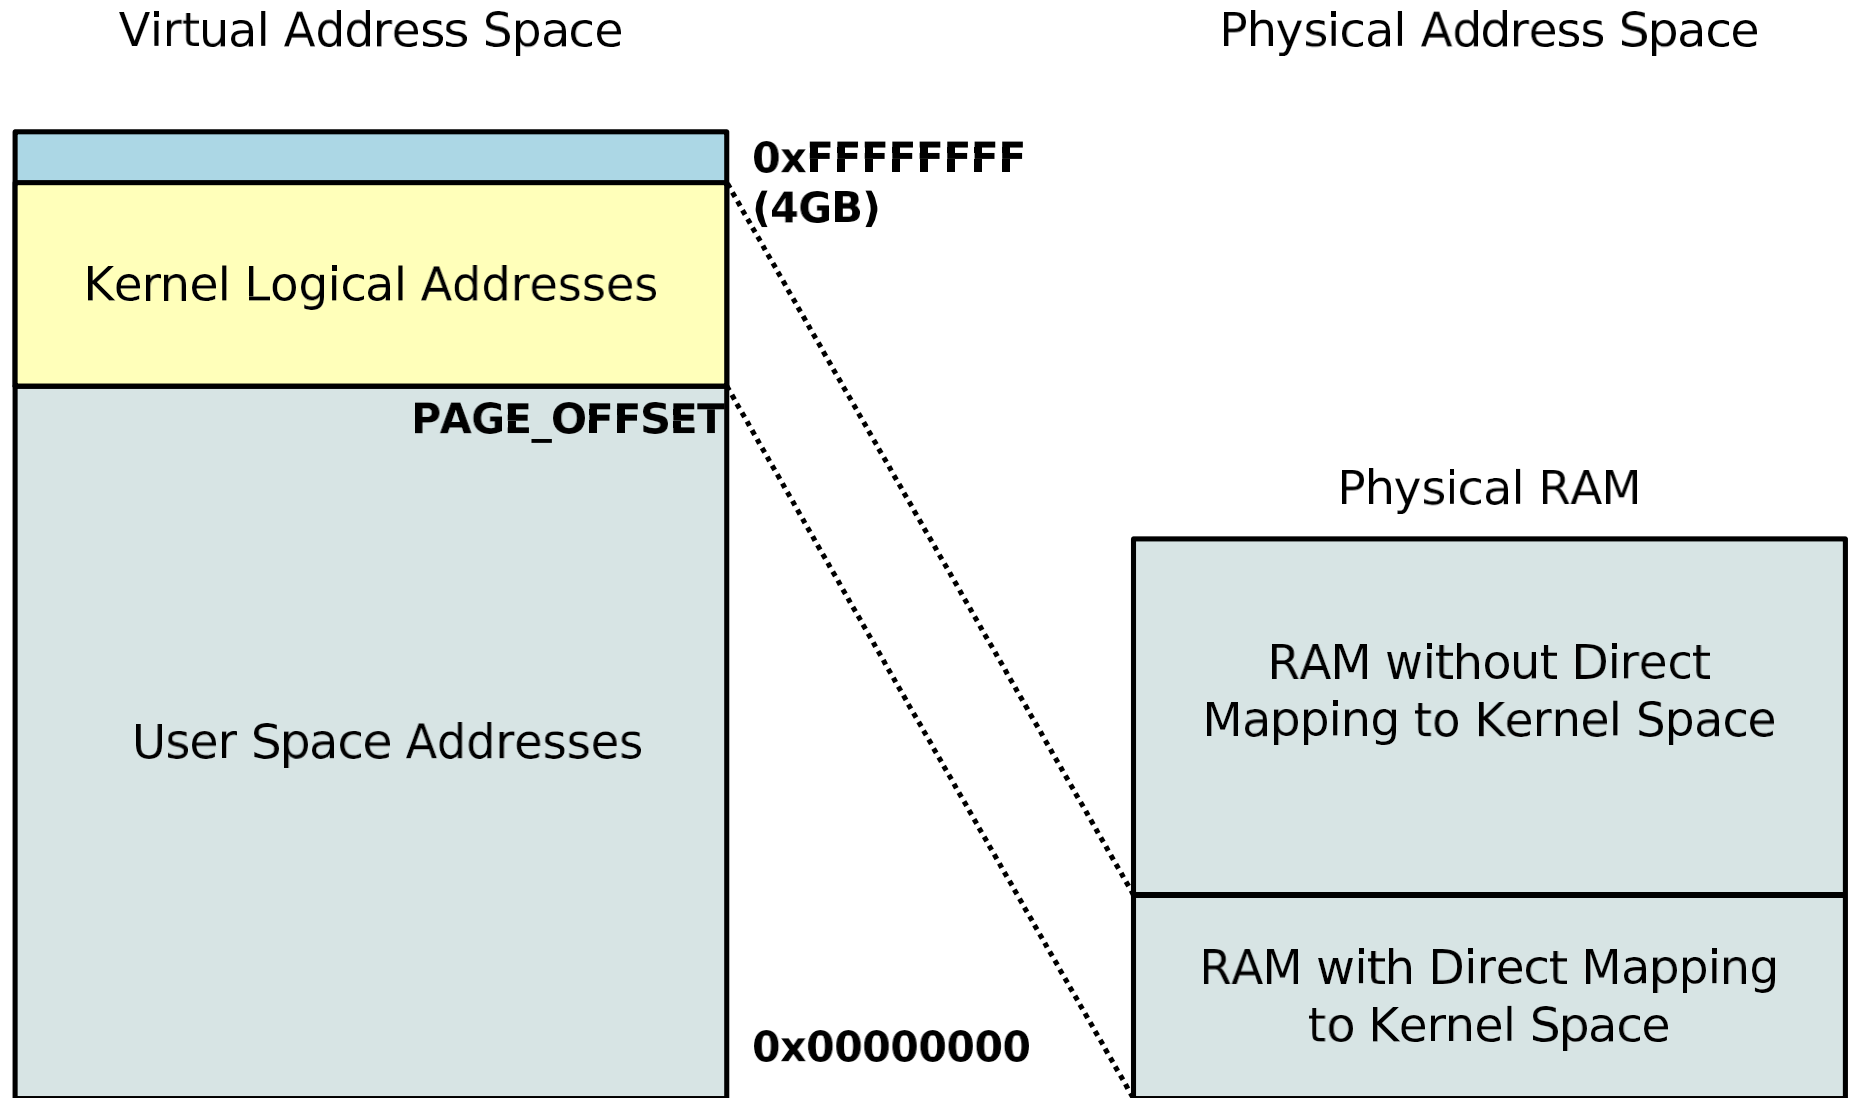
- How kernel manage physical RAM?
  - Such as allocate a physical frame to user process
- Must map physical RAM to kernel AS
  - For 32-bit, kernel address space (**AS**) is 1GB
    - No problem to handle small RAM ( $\leq 896\text{MB}$ )
    - How about **large** physical RAM  $\geq 1\text{GB}$ 
      - Even **larger** than the kernel AS

# Large physical RAM

---

- For large memory systems (more than ~1GB RAM), **not** all of the physical RAM can be **linearly** mapped into the kernel's address space.
  - Kernel address space is the top 1GB of virtual address space, by default.
  - Further, 128 MB is reserved at the top of the kernel's memory space for non-contiguous allocations
    - See `vmalloc()` described later

# Kernel Logical Addresses (Large Mem)



# Large physical RAM

---

- As a result, only the bottom part of physical RAM has a kernel logical address
  - Note that on 64-bit systems, kernel AS is much larger than physical RAM, this case never happens
    - 39-bit kernel AS is 512GB
    - 48-bit kernel AS is 256TB
- Thus, in a large memory situation, only the bottom part of physical RAM is mapped linearly into kernel logical address space
  - Top 128MB guarantees that there is always enough kernel address space to accommodate all the RAM



# Kernel Virtual Addresses

---

- Kernel Virtual Addresses are addresses in the region above the kernel logical address mapping.
- Kernel Virtual Addresses are used for non-contiguous memory mappings
  - Often for large buffers which could potentially be unable to get physically contiguous regions allocated.
- Also referred to as the `vmalloc()` area

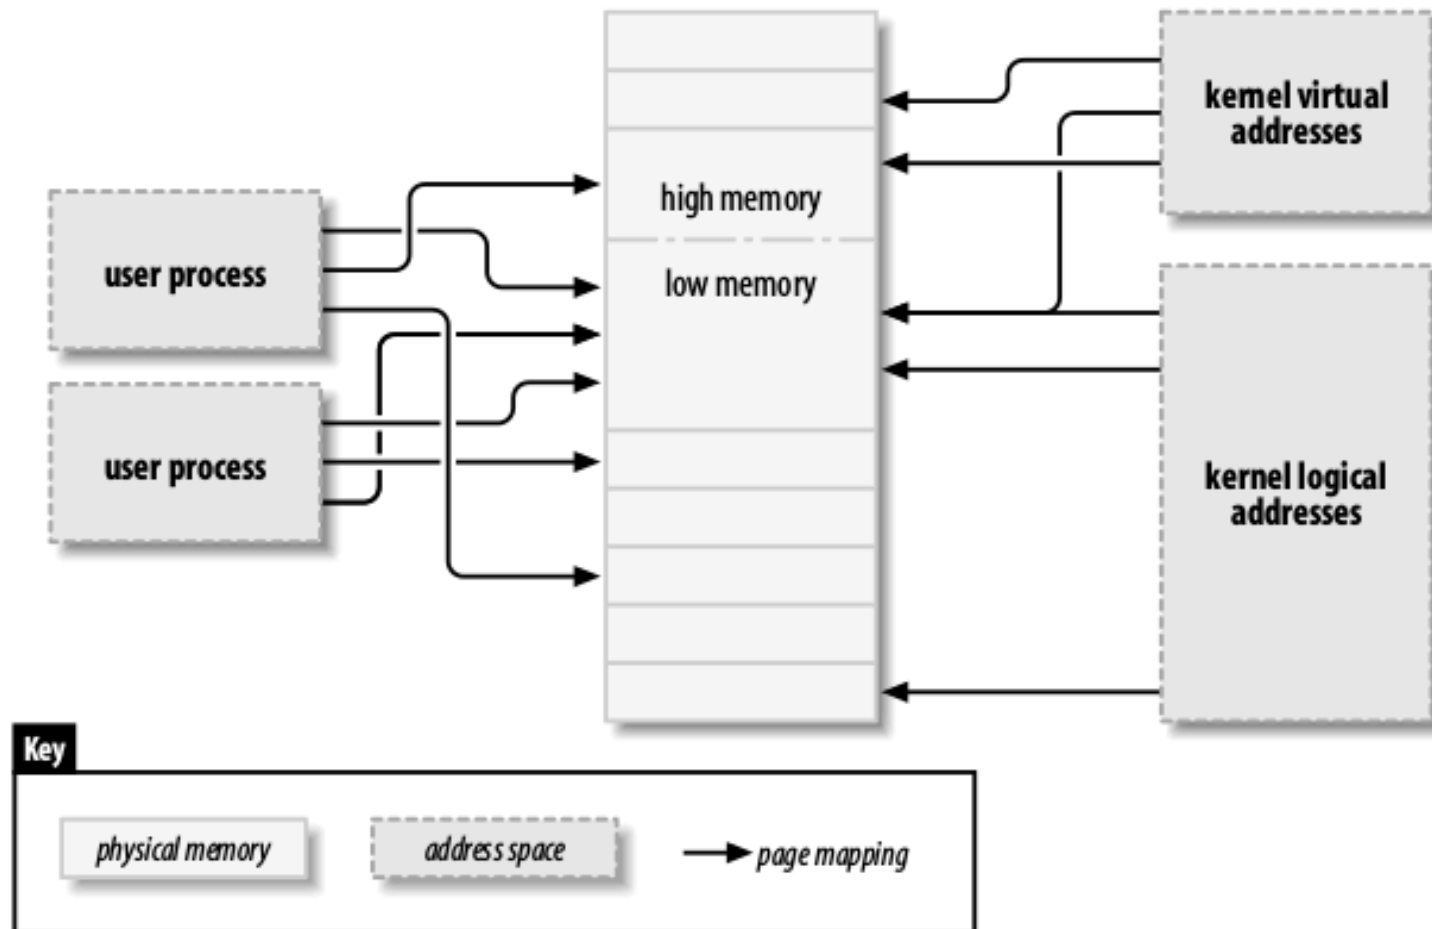
# Virtual Addresses - Summary

---

- There are three kinds of virtual addresses in Linux.
  - The terminology varies, even in the kernel source, but the definitions in *Linux Device Drivers, 3<sup>rd</sup> Edition*, chapter 15, are somewhat standard.
  - LDD 3 can be downloaded for free at:  
<https://lwn.net/Kernel/LDD3/>

# Virtual Addresses - Linux

- User virtual address
- Kernel virtual address
- Kernel logical address



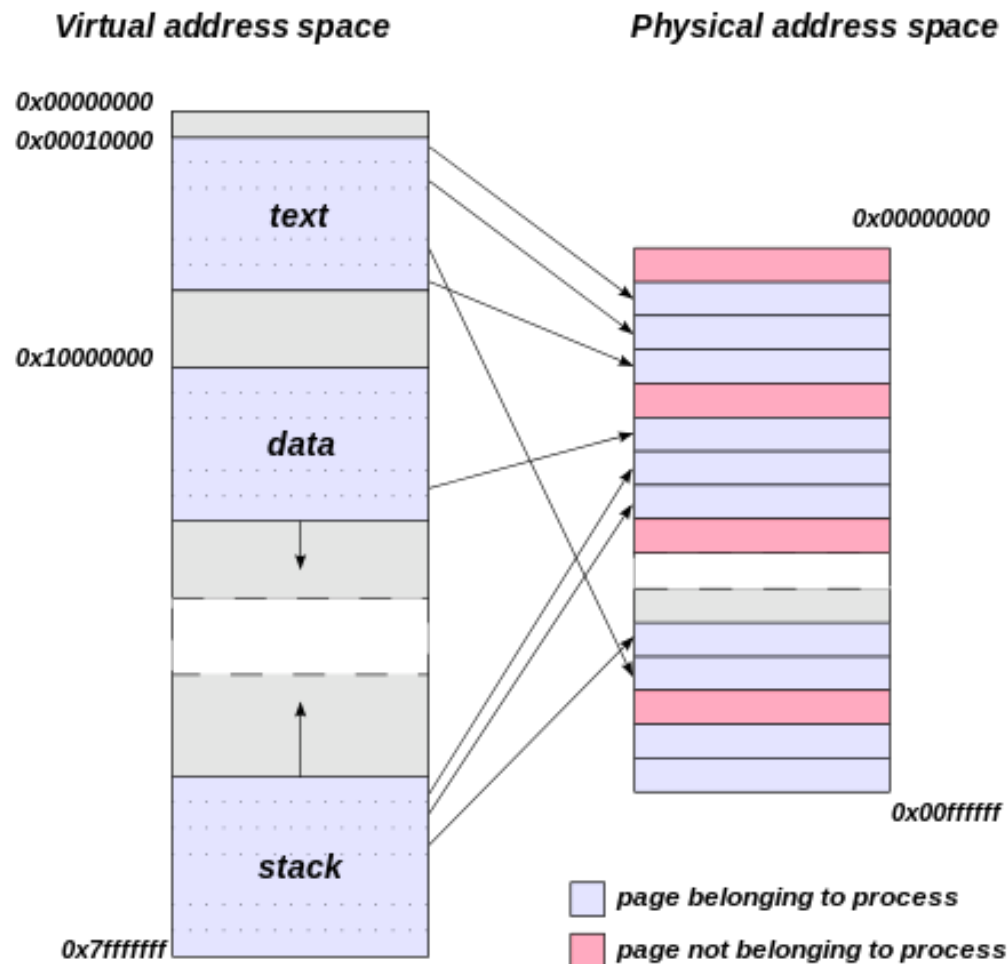
# Virtual Addresses - Linux

---

- User virtual address, Kernel virtual address, Kernel logical address
  - All belong to virtual address, but has different mappings to physical memory
  - Who has contiguous mapping?
  - Who has non-contiguous mapping?
- All these mapping can be done by MMU
  - By walking page table
  - Review: what is page table?

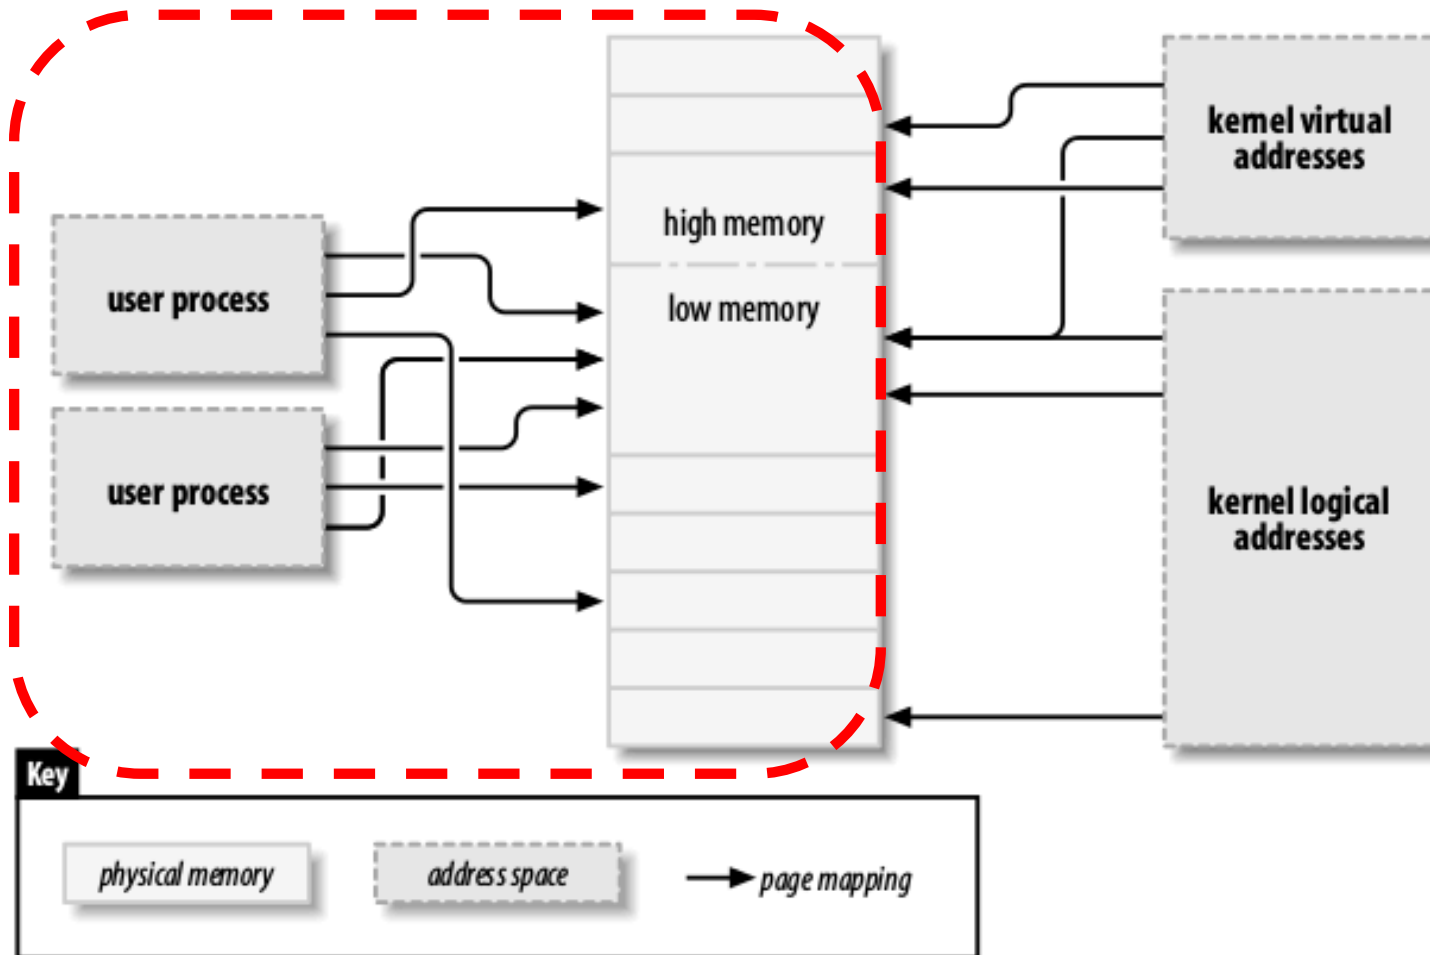
# User space page table mapping

- Non-contiguous



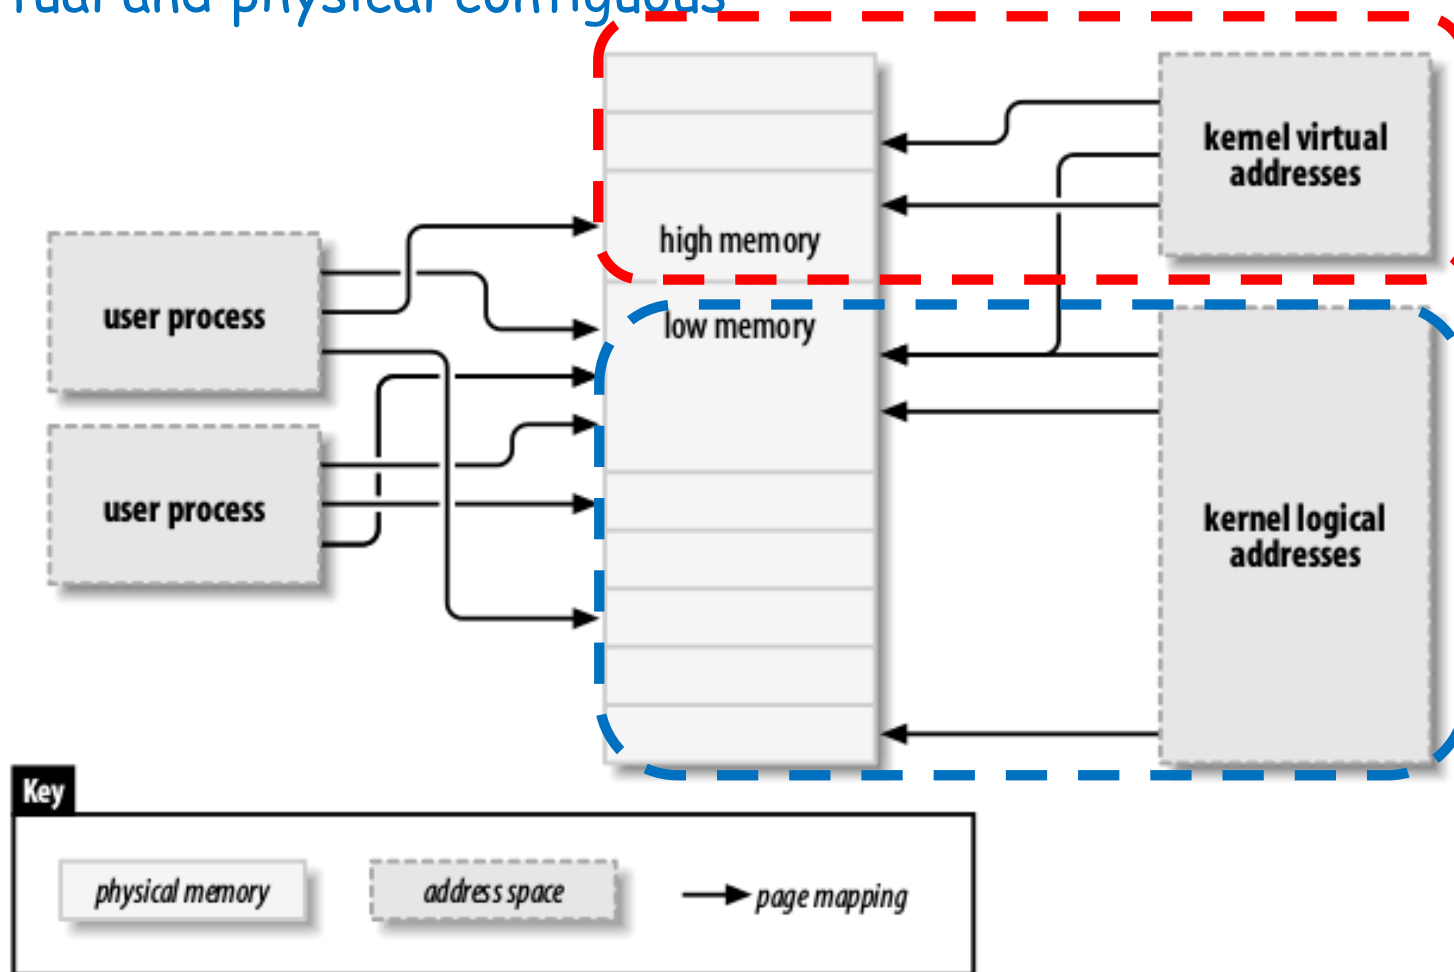
# User space page table mapping

- Virtual contiguous, physical non-contiguous



# Kernel space page table mapping

- Virtual contiguous, physical non-contiguous
- Virtual and physical contiguous




# Basic Page Table Mappings

---

User Virtual Address Space

Page Table

Physical Address Space



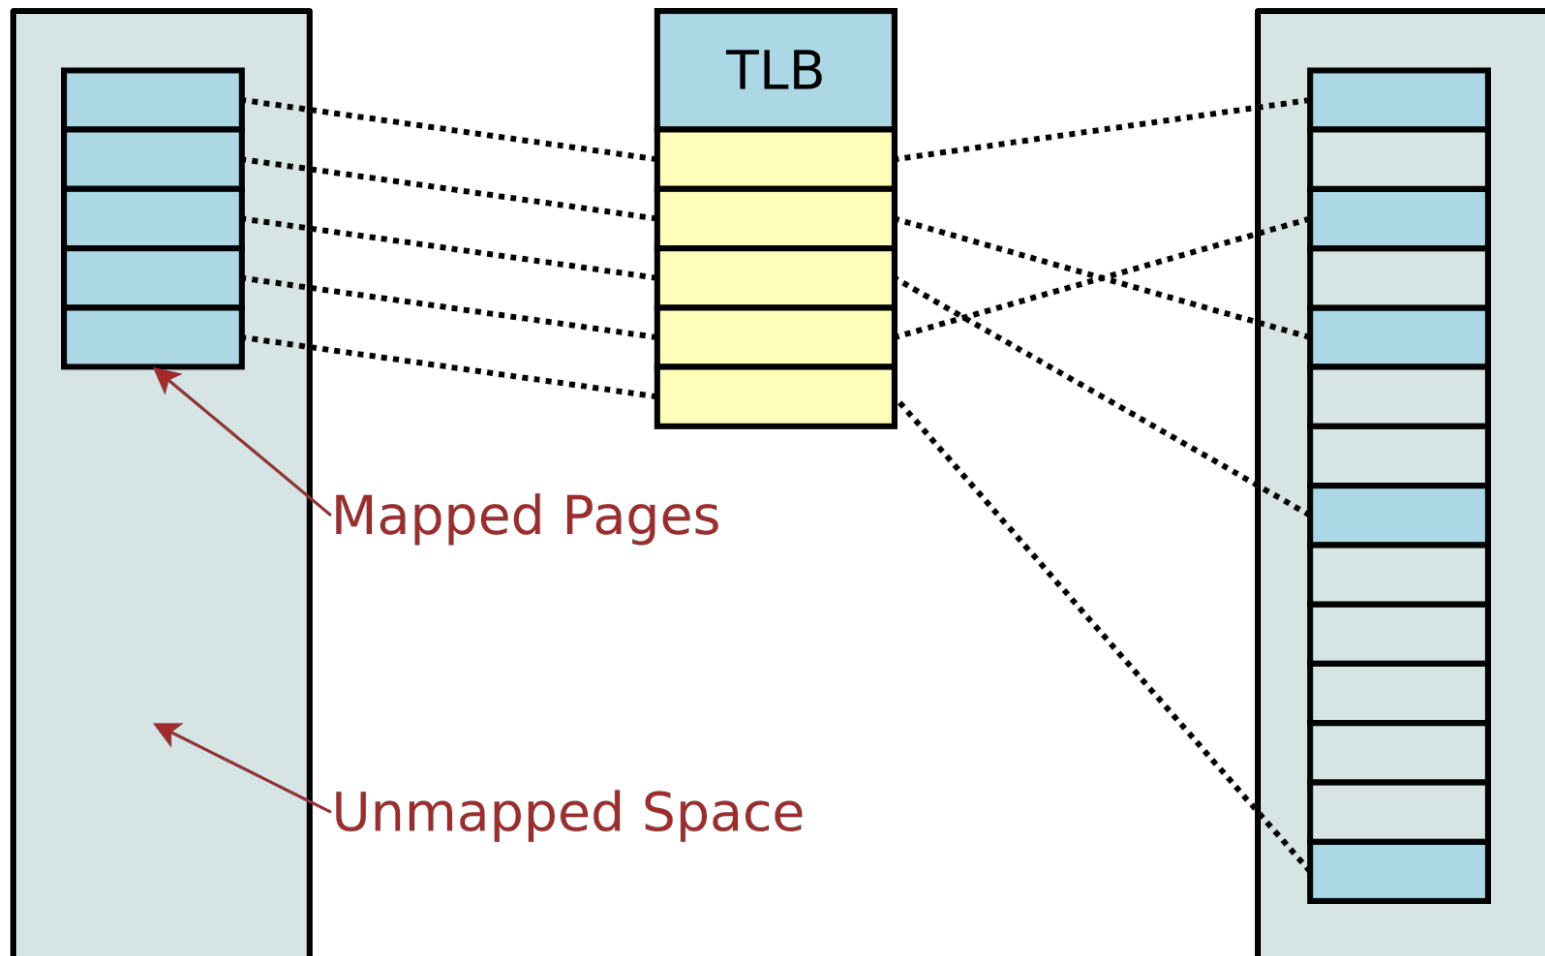


# Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space




# Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space

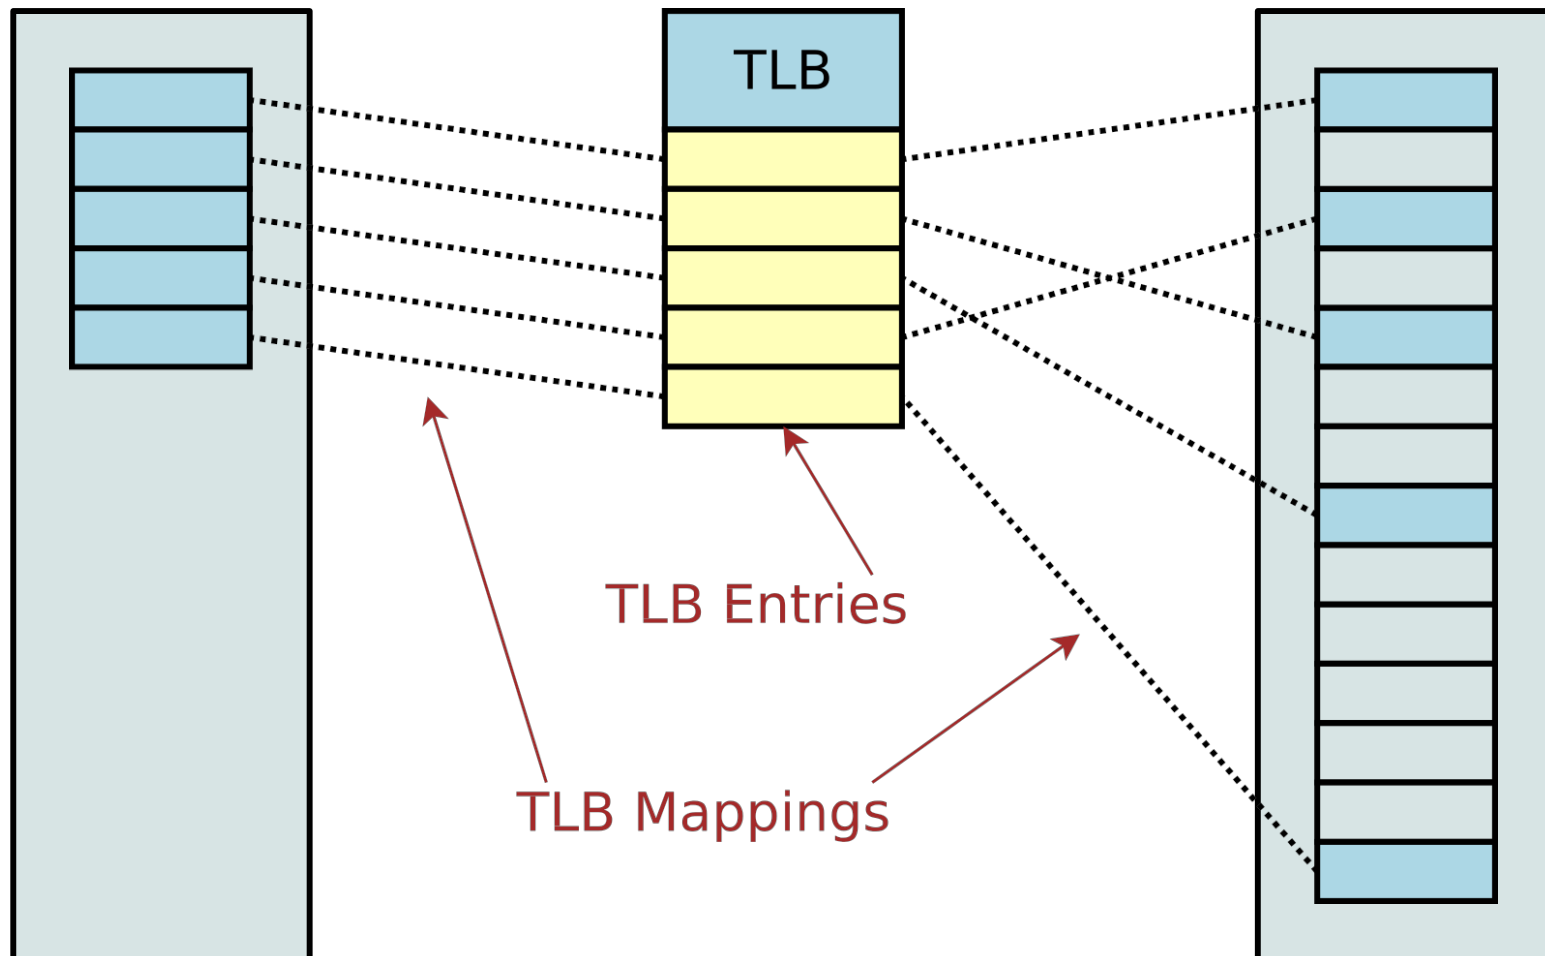


# Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space




# Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space



# Page Table Mappings

---

- Mappings to virtually contiguous regions do not have to be physically contiguous.
  - This makes memory easier to allocate.
  - Almost all user space code does not need physically contiguous memory.

# Multiple Processes

---

- Each process has its own mapping.
  - The same virtual addresses in different processes point to different physical addresses in other processes


# Multiple Processes - Process 1

---

User Virtual Address Space

Page Table

Physical Address Space




# Multiple Processes - Process 2

---

User Virtual Address Space

Page Table

Physical Address Space





# Shared Memory

---

- Shared memory is easily implemented with an MMU.
  - Simply map the same physical frame into two different processes.
  - The virtual addresses need not be the same.
    - If pointers to values inside a shared memory region are used, it might be important for them to have the same virtual addresses, though.

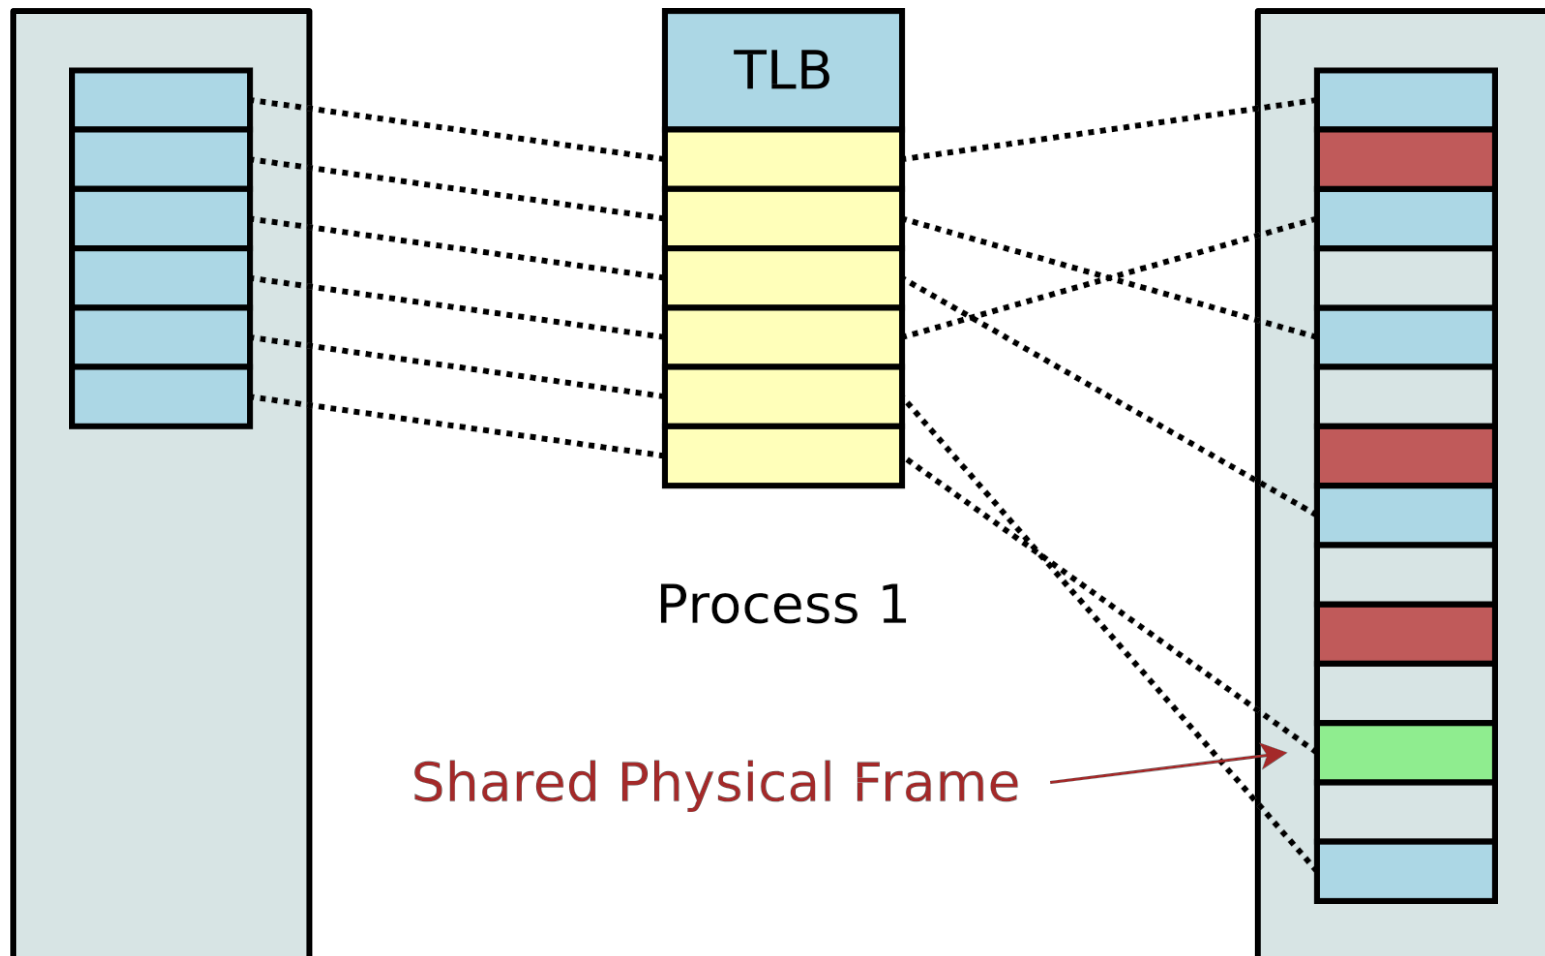
# Shared Memory - Process 1

---

User Virtual Address Space

Page Table

Physical Address Space

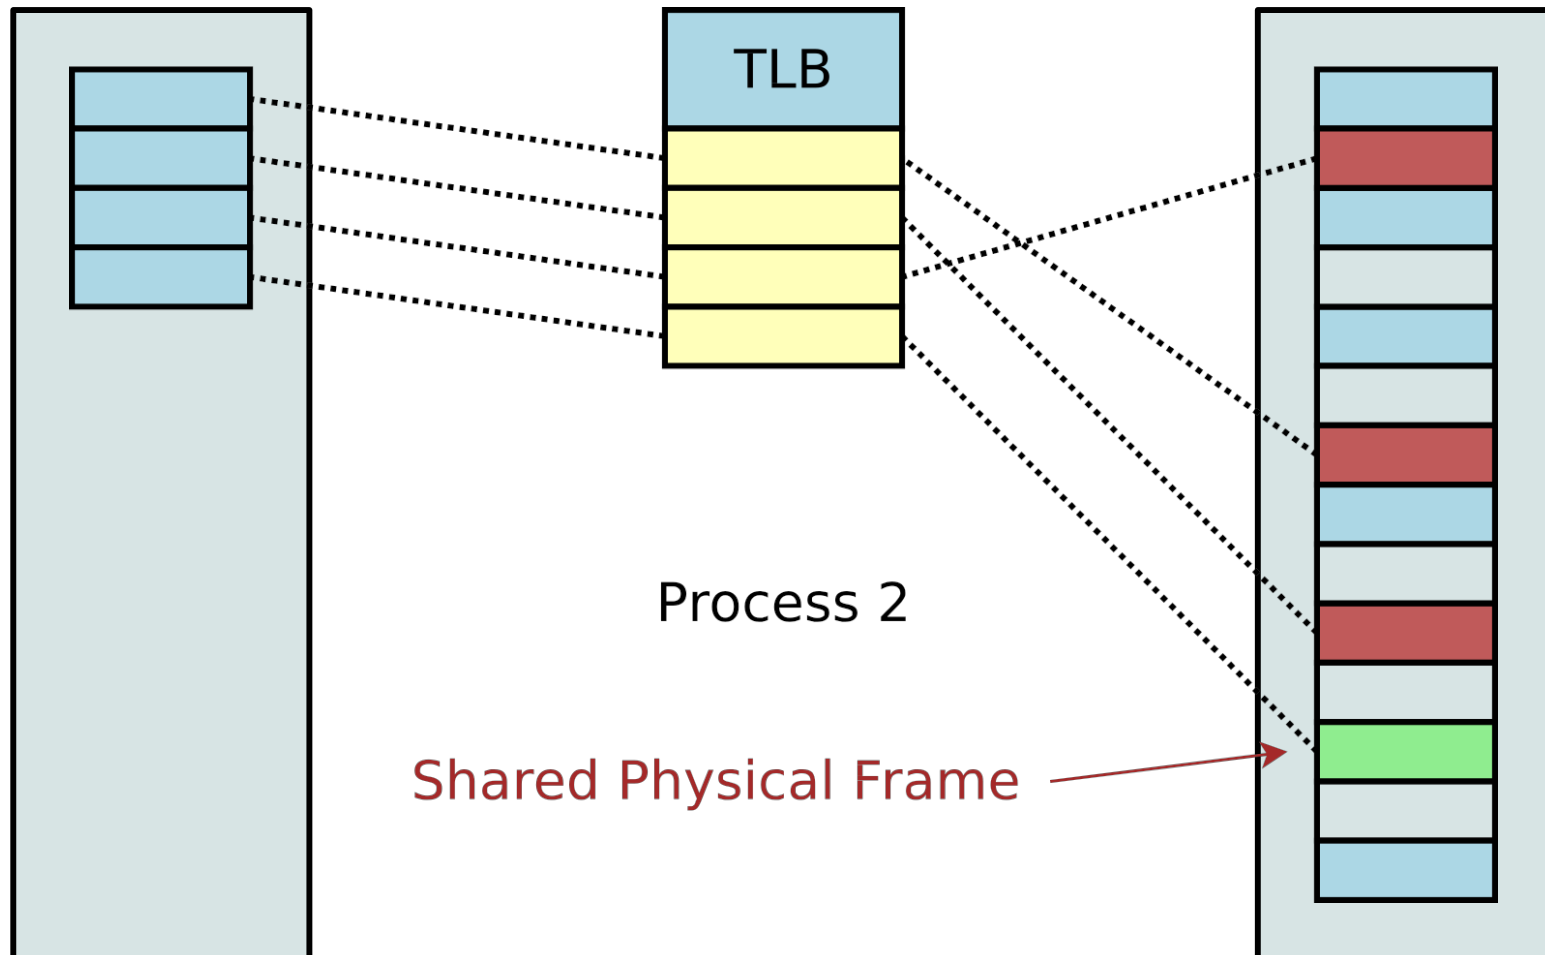


# Shared Memory - Process 2

User Virtual Address Space

Page Table

Physical Address Space



# Shared Memory

---

- Note in the previous example, the shared memory region was mapped to different virtual addresses in each process.
- The `mmap()` system call allows the user space process to request a virtual address to map the shared memory region.
  - The kernel may not be able to grant a mapping at this address, causing `mmap()` to return failure.

# Kernel Logical Addresses - Repeat

---

- Kernel Logical addresses can be converted to and from physical addresses using the macros:

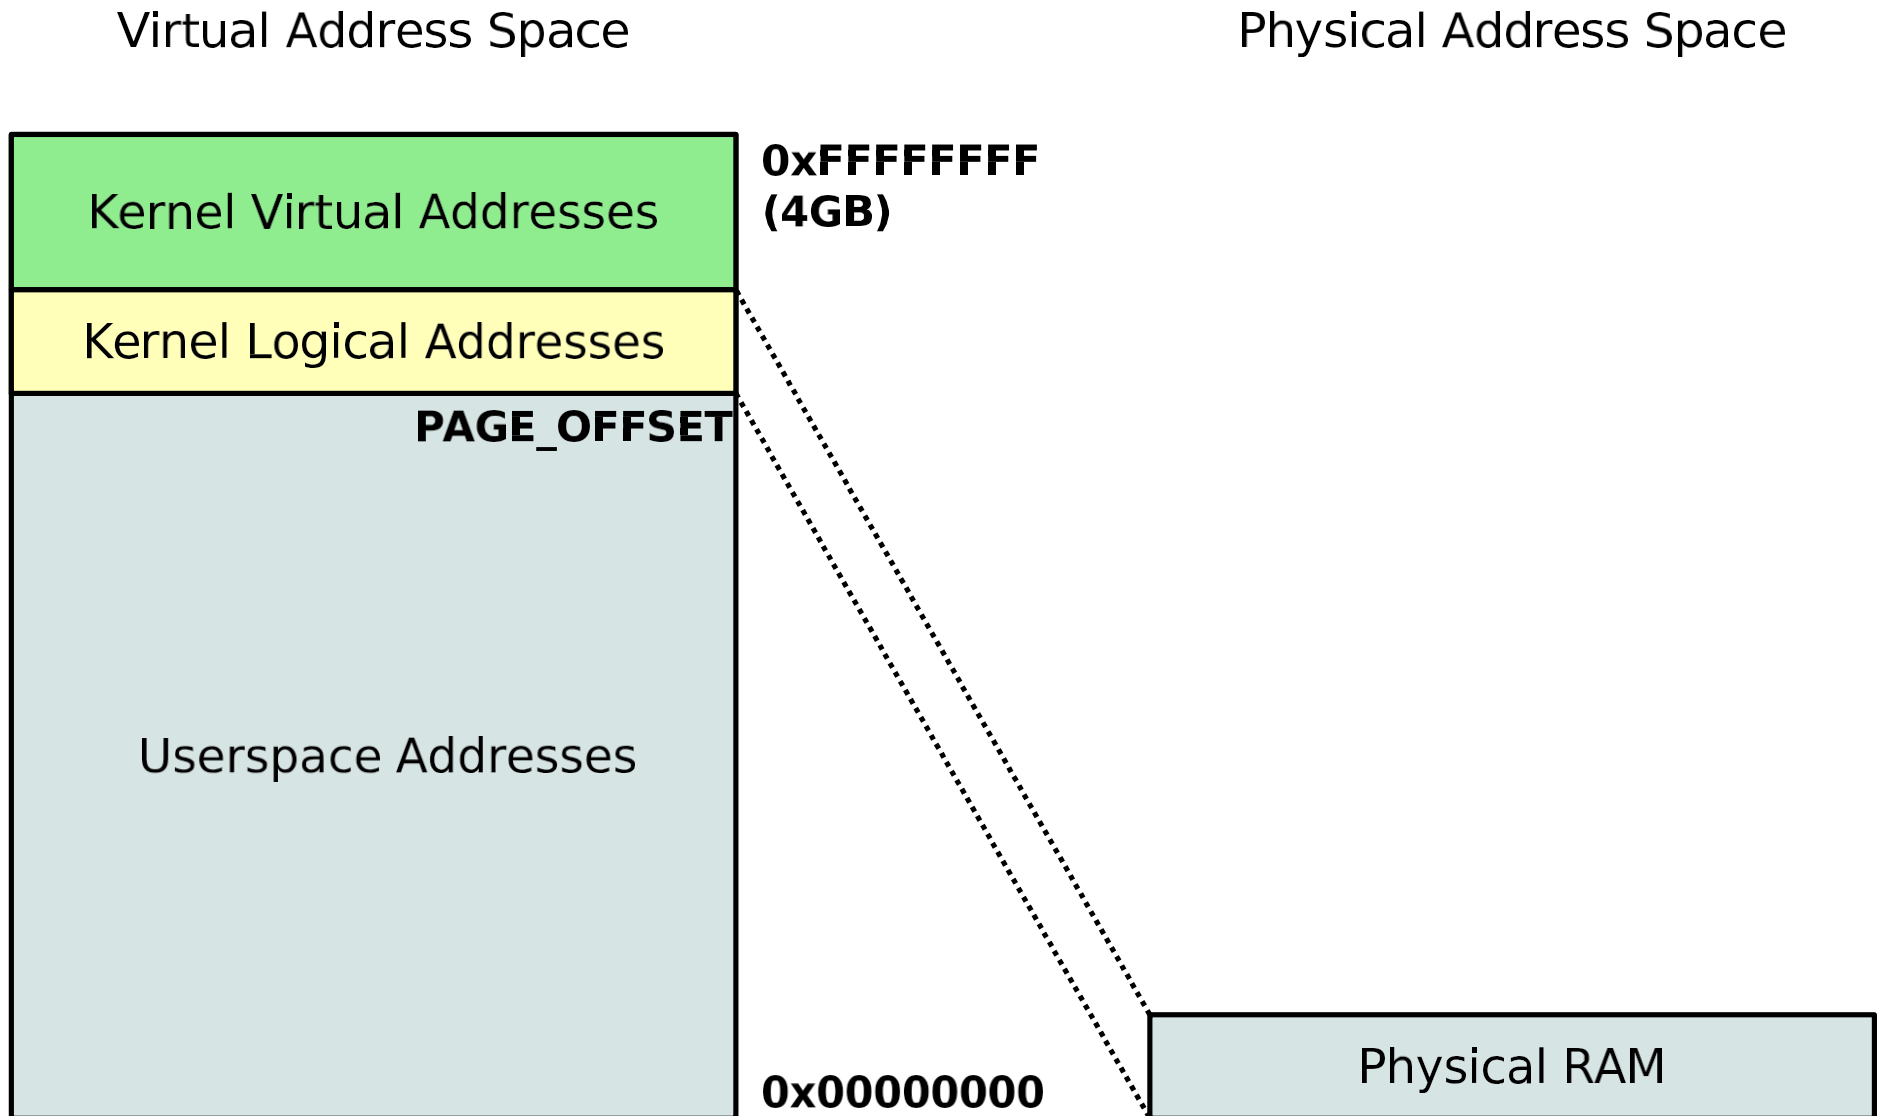
    \_\_pa(x)

    \_\_va(x)

- For low-memory systems (below ~1G of RAM) Kernel Logical address space starts at PAGE\_OFFSET and goes through the end of *physical* memory.

# Kernel Virtual Addresses

---



# Kernel Virtual Addresses

---

- In the small memory model, as shown, since all of RAM can be represented by logical addresses, all virtual addresses will also have logical addresses.
  - One mapping in virtual address area
  - One mapping in logical address area

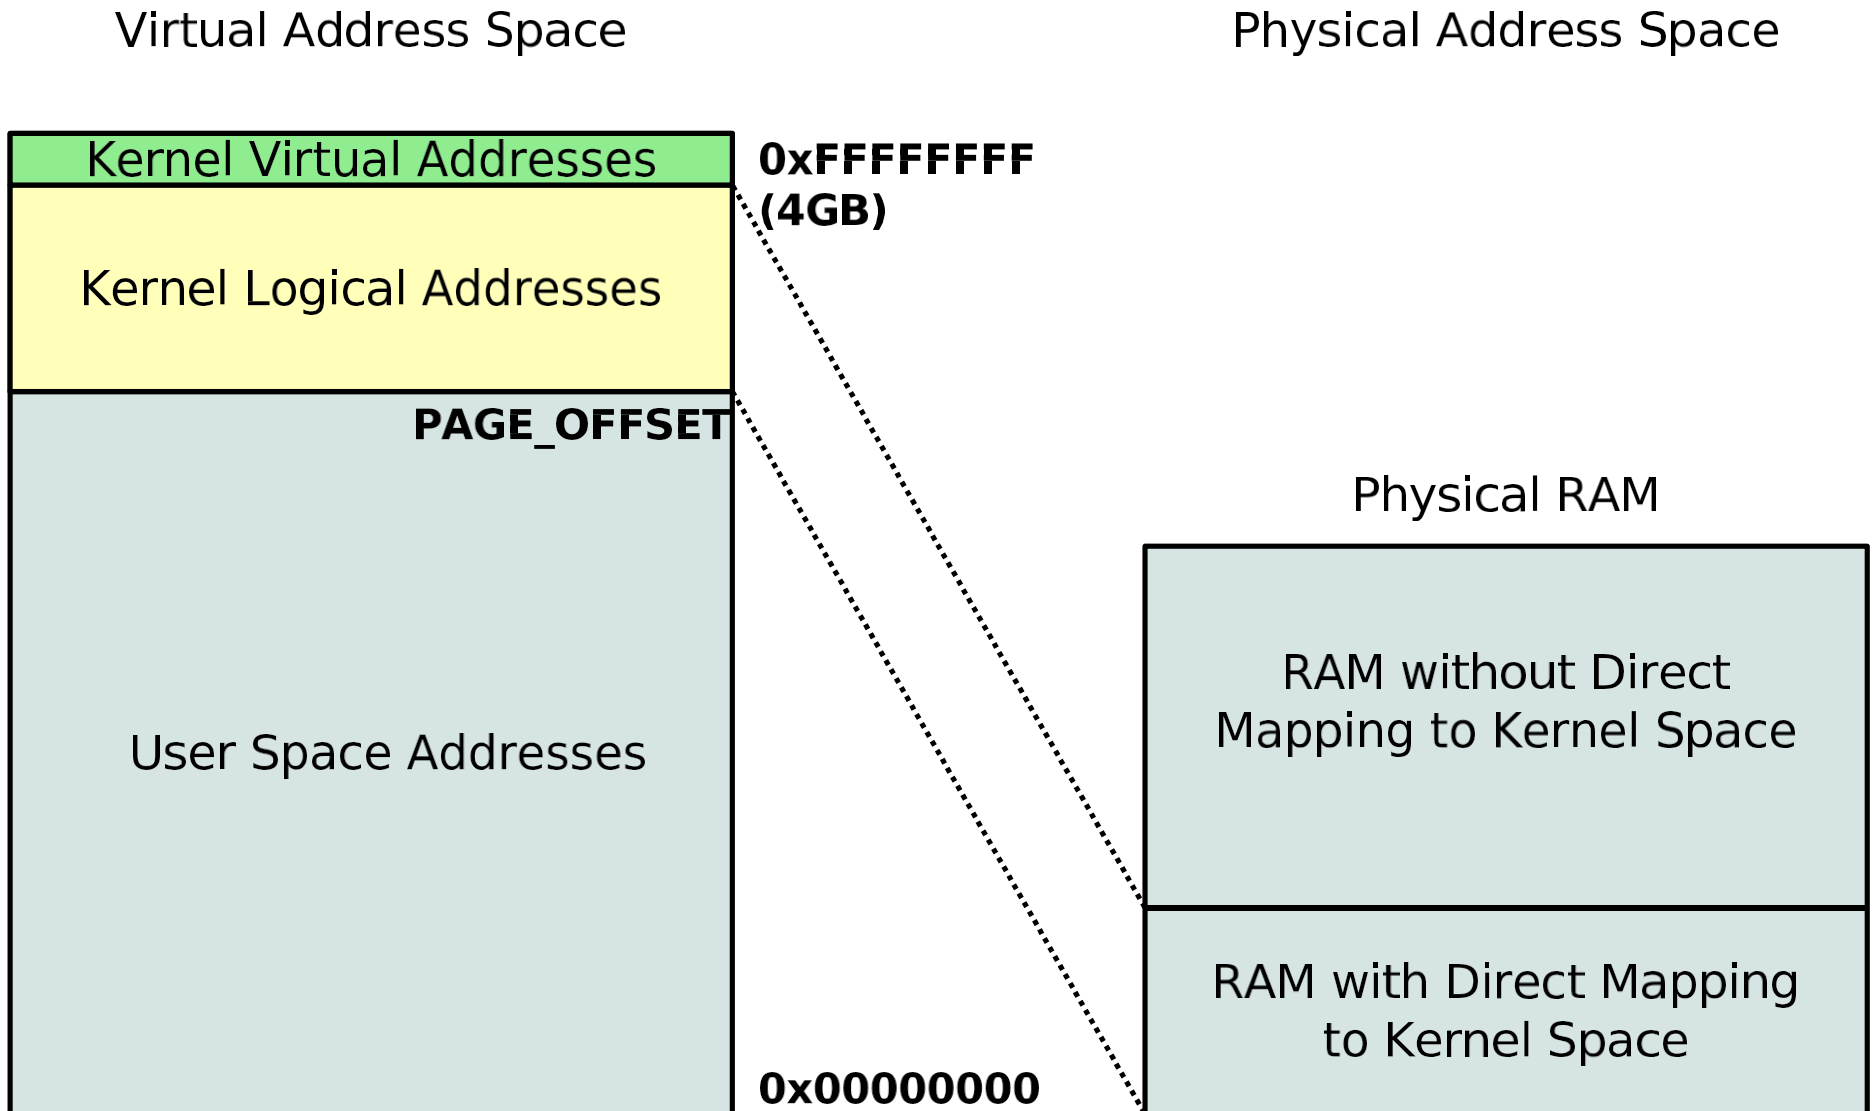
# Kernel Virtual Addresses

---

- The important difference is that memory in the kernel virtual address area (or `vmalloc()` area) is non-contiguous physically.
  - This makes it easier to allocate, especially for large buffers
  - This makes it unsuitable for DMA



# Kernel Virtual Addresses (Large Mem )



# Kernel Virtual Addresses

---

- In a large memory situation, the kernel virtual address space is smaller, because there is more physical memory.
  - An interesting case, where more memory means less virtual address space.
  - In 64-bit, of course, this doesn't happen, as `PAGE_OFFSET` is large, and there is much more virtual address space.

# User Virtual Addresses

---

- User Virtual Addresses represent memory used by user space programs.
  - This is most of the memory on most systems
  - This is where most of the complication is
- User virtual addresses are all addresses below `PAGE_OFFSET`.
- Each process has its own mapping

# User Virtual Addresses

---

- Unlike kernel logical addresses, which use a linear mapping between virtual and physical addresses, user space processes make full use of the MMU.
  - Only the used portions of RAM are mapped
  - Memory is not contiguous
  - Memory may be swapped out
  - Memory can be moved

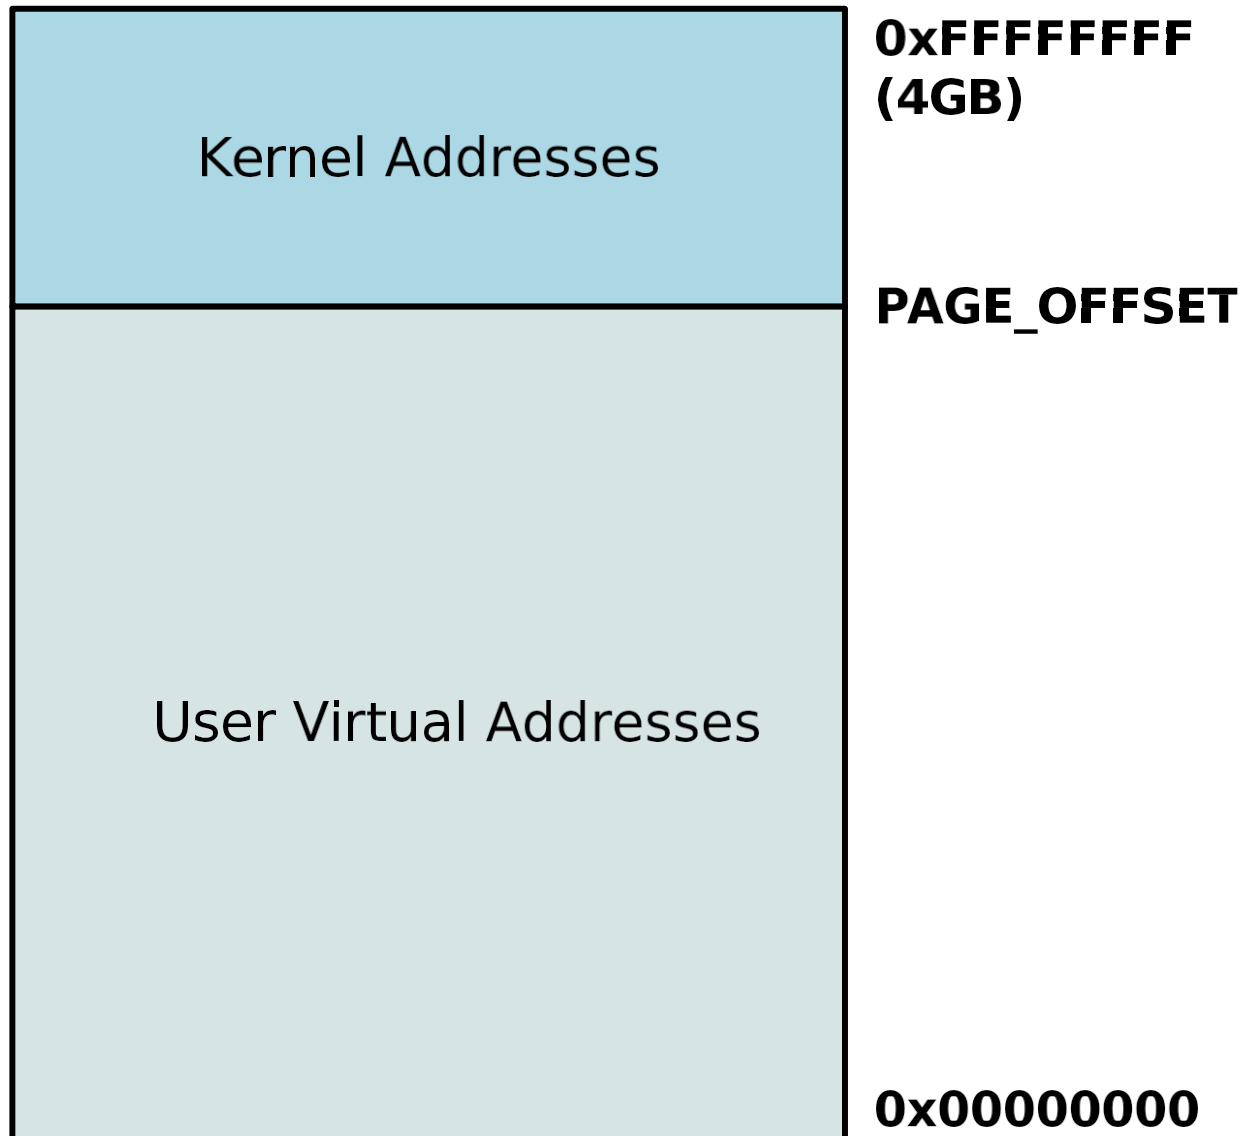
# User Virtual Addresses

---

- Since user virtual addresses are not guaranteed to be swapped in, or even allocated at all, user pointers are not suitable for use with kernel buffers or DMA, by default.
- Each process has its own memory map
  - `mm_struct`
- At context switch time, the memory map of the new process will be switched in
  - This is part of the context switch overhead

# User Virtual Addresses

---



- Each process will have its own mapping for user virtual addresses
- The mapping is changed during context switch


# Show me the code

---

- Talk is cheap, show me the code

---

# Memory Management Unit





# The MMU

---

- The Memory Management Unit (MMU) is a hardware component which manages virtual address mappings
  - Maps virtual addresses to physical addresses
- The MMU operates on basic units of memory called pages
  - Page size varies by architecture
  - Some architectures have configurable page sizes

# The MMU

---

- Common page sizes:
  - ARM – 4k
  - ARM64 – 4k or 64k
  - MIPS – Widely Configurable
  - x86 – 4k
    - *Architectures which are configurable are configured at kernel build time.*

# The MMU

---

- Terminology
  - A **page** is a unit of memory sized and aligned at the page size.
  - A **page frame**, or frame, refers to a page-sized and page-aligned physical memory block.
    - *A page is somewhat abstract, where a frame is concrete*
    - *In the kernel, the abbreviation **pfn**, for **page frame number**, is often used to refer to refer to physical page frames*

# The MMU

---

- The MMU operates in pages
  - The MMU maps physical frames to virtual addresses.
  - The TLB holds the entries of the mapping
    - Virtual address
    - Physical address
    - Permissions
  - A memory map for a process will contain many mappings

# Page Faults

---

- When a process accesses a region of memory that is not mapped, the MMU will generate a page fault exception.
  - The kernel handles page fault exceptions regularly as part of its memory management design.

# Lazy Allocation

---

- The kernel will not allocate pages requested by a process immediately.
  - The kernel will wait until those pages are actually used.
  - This is called lazy allocation and is a performance optimization.
    - For memory that doesn't get used, allocation never has to happen!

# Lazy Allocation

---

- Process
  - When memory is requested, the kernel simply creates a record of the request, and then returns (quickly) to the process, without updating the TLB.
  - When that newly-allocated memory is touched, the CPU will generate a page fault, because the CPU doesn't know about the mapping

# Lazy Allocation

---

- Process (cont)
  - In the page fault handler, the kernel determines that the mapping is valid (from the kernel's point of view).
  - The kernel updates the TLB with the new mapping
  - The kernel returns from the exception handler and the user space program resumes.



# Lazy Allocation

---

- In a lazy allocation case, the user space program never is aware that the page fault happened.
  - The page fault can only be detected in the time that was lost to handle it.
- For processes that are time-sensitive, data can be pre-faulted, or simply touched, at the start of execution.
  - Also see `mlock()` and `mlockall()` for pre-faulting.

# Page Tables

---

- The entries in the TLB are a limited resource.
- Far more mappings can be made than can exist in the TLB at one time.
- The kernel must keep track of all of the mappings all of the time.
- The kernel stores all this information in the page tables.

# Page Tables

---

- Since the TLB can only hold a limited subset of the total mappings for a process, some mappings will not have TLB entries.
  - When these addresses are touched, the CPU will generate a page fault, because the CPU has knowledge of the mapping.

# Page Tables

---

- When the page fault handler executes in this case, it will:
  - Find the appropriate mapping for the offending address in the kernel's page tables
  - Select and remove an existing TLB entry
  - Create a TLB entry for the page containing the address
  - Return to the user space process
    - *Observe the similarities to lazy allocation handling*

# Swapping

---

- When memory allocation is high, the kernel may swap some frames to disk to free up RAM.
  - Having an MMU makes this possible.
- The kernel can copy a frame to disk and remove its TLB entry.
- The frame can be re-used by another process

# Swapping

---

- When the frame is needed again, the CPU will generate a page fault (because the address is not in the TLB)
- The kernel can then, at page fault time:
  - Put the process to sleep
  - Copy the frame from the disk into an unused frame in RAM
  - Fix the page table entry   Wake the
  - process

# Swapping

---

- Note that when the page is restored to RAM, it's not necessarily restored to the same physical frame where it originally was located.
- The MMU will use the same virtual address though, so the user space program will not know the difference
  - *This is why user space memory cannot typically be used for DMA.*

# Takeaway

---

- Virtual memory in Linux
  - Kernel logical address
  - Kernel virtual address
  - User virtual address
- Contiguous vs non-contiguous
  - Virtually, physically
- Linux data struct
  - `mm_struct`, `switch_mm`, `mm_struct.pgd`
- Lazy allocation
  - Page fault



# Page table quiz

---

- 1) In 32-bit architecture, for 1-level page table, how large is the whole page table?
- 2) In 32-bit architecture, for 2-level page table, how large is the whole page table?
  - 1) How large for the 1<sup>st</sup> level PGT?
  - 2) How large for the 2<sup>nd</sup> level PGT?
- 3) Why can 2-level PGT save memory?
- 4) 2-level page table walk example
  - 1) Page table base register holds 0x0061 9000
  - 2) Virtual address is 0xf201 5202
  - 3) Page table base register holds 0x1051 4000
  - 4) Virtual address is 0x2190 7010